



# Computer Networks

CMSC 417 : Spring 2024



COMPUTER SCIENCE  
UNIVERSITY OF MARYLAND

## Transport Layer Protocols (UDP, TCP) (Textbook chapter 5)

Nirupam Roy

Tu-Th 2:00-3:15pm

CSI 2117

March 7<sup>th</sup>, 2024



## Connection-less mux/demux:

- User Datagram Protocol (UDP)

## Connection-oriented mux/demux:

- Transmission Control Protocol (TCP)

# Connectionless demultiplexing

- Create sockets with port numbers:

```
DatagramSocket mySocket1 = new  
    DatagramSocket(99111);
```

```
DatagramSocket mySocket2 = new  
    DatagramSocket(99222);
```


- UDP socket identified by two-tuple:

(dest IP address, dest port number)

- When host receives UDP segment:
  - checks destination port number in segment
  - directs UDP segment to socket with that port number
- IP datagrams with different source IP addresses and/or source port numbers directed to same socket

# Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



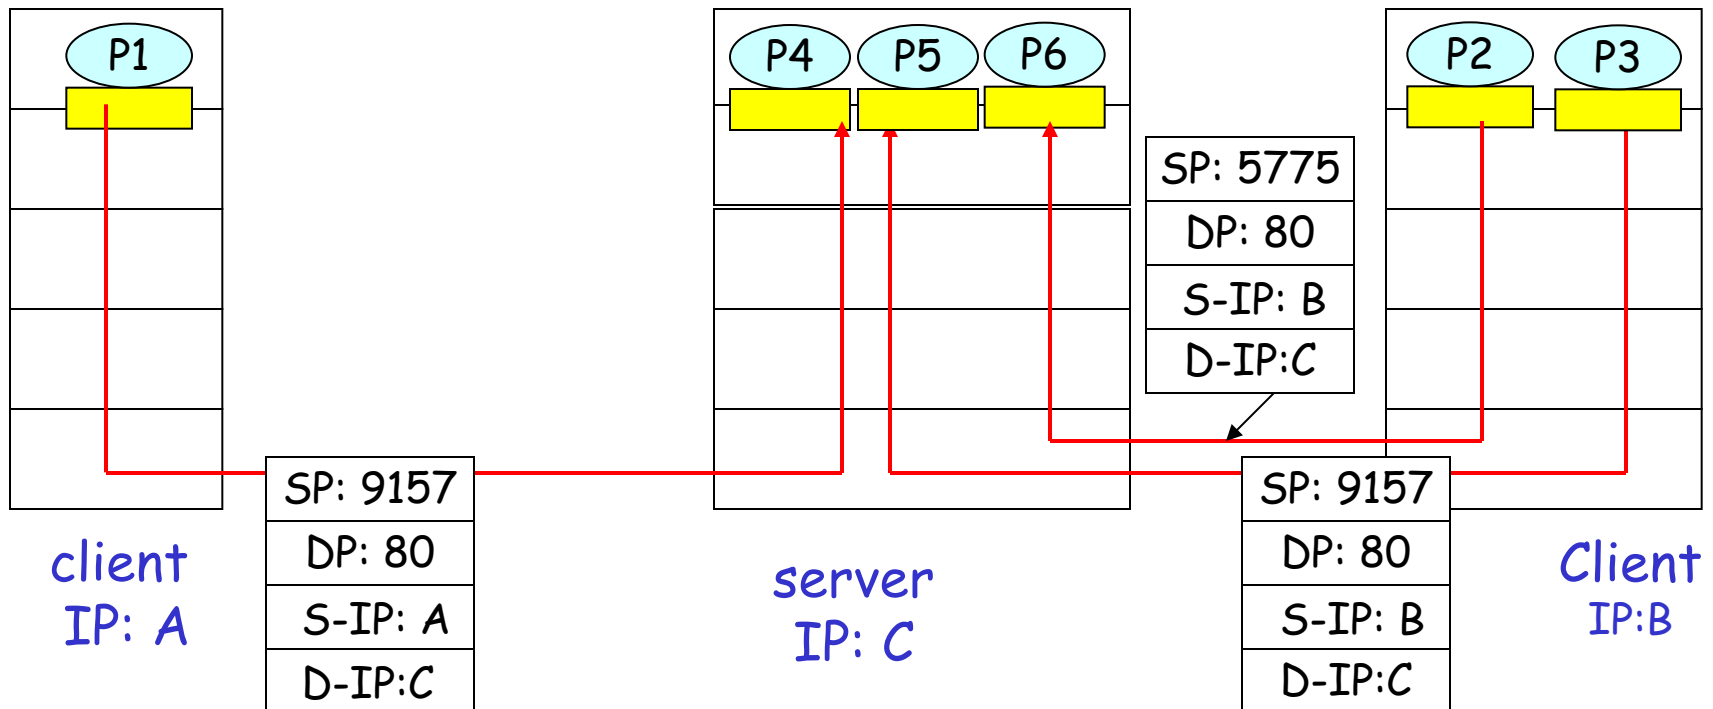
SP provides "return address"

# Connection-oriented demux

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- recv host uses all four values to direct segment to appropriate socket
- Server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

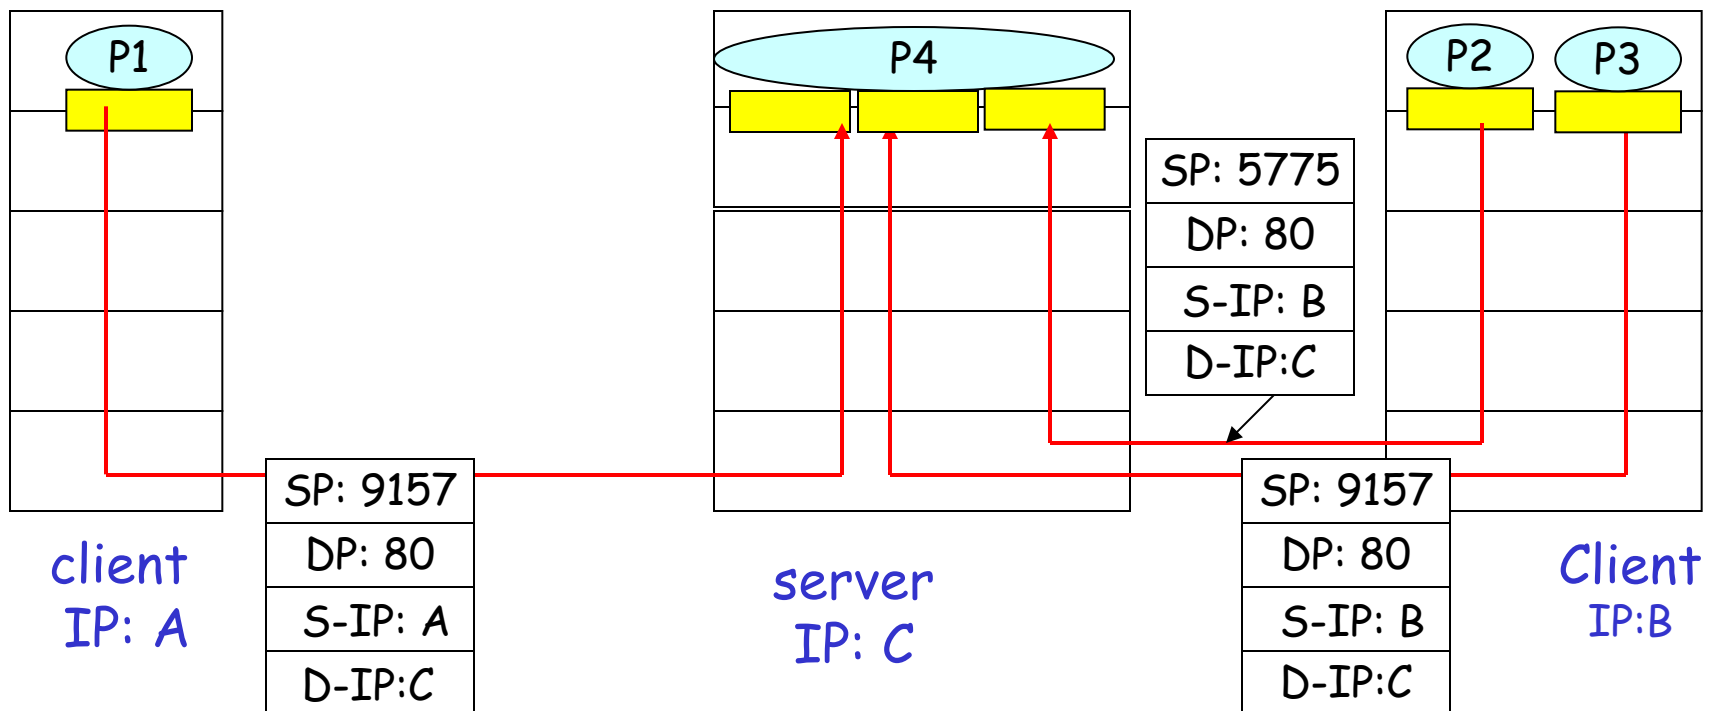
# Connection-oriented demux (cont)

 = socket       = process



# Connection-oriented demux: Threaded Web Server

 = socket       = process



## (2) Abstraction of services



# IP Protocol Stack: Key Abstractions

<b>Application</b>	Applications	
<b>Transport</b>	Streams of data	Messages
<b>Network</b>	Best-effort <i>global</i> packet delivery	
<b>Link</b>	Best-effort <i>local</i> packet delivery	

- Transport layer is where we “pay the piper”
  - Provide applications with good abstractions
  - Without support or feedback from the network

# End-to-end Protocols (UDP, TCP)

# End-to-end Protocols

- Common properties that a transport protocol can be expected to provide
  - **Guarantees** message delivery
  - Delivers messages in the **same order** they were sent
  - Delivers at most **one copy** of each message
  - Supports **arbitrarily large** messages
  - Supports synchronization between the sender and the receiver
  - Allows the receiver to apply flow control to the sender
  - Supports multiple application processes (Mux/DeMux)

# End-to-end Protocols

- Typical limitations of the network on which transport protocol will operate
  - Drop messages/Packet loss
  - Reorder messages/Out of order delivery
  - Deliver duplicate copies of a given message
  - Limit messages to some finite size
  - Deliver messages after an arbitrarily long delay

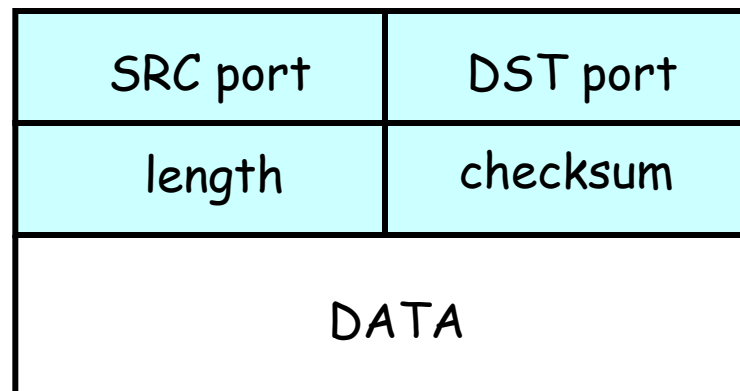
# User Datagram Protocol (UDP)

# UDP: User Datagram Protocol [RFC 768]

- ❑ Bare minimum Internet transport protocol
- ❑ “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- ❑ *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

# User Datagram Protocol (UDP)

- **Datagram messaging service**
  - Demultiplexing: port numbers
  - Detecting corruption: checksum
- **Lightweight communication between processes**
  - Send and receive messages
  - Avoid overhead of any ordered, reliable delivery



# UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

## Sender:

- ❑ treat segment contents as sequence of 16-bit integers
- ❑ checksum: addition (1’s complement sum) of segment contents
- ❑ sender puts checksum value into UDP checksum field

## Receiver:

- ❑ compute checksum of received segment
- ❑ check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.

Is UDP service == IP service ??




# Advantages of UDP

- **Fine-grain control**
  - UDP sends as soon as the application writes
- **No connection set-up delay**
  - UDP sends without establishing a connection
- **No connection state**
  - No buffers, parameters, sequence #s, etc.
- **Small header overhead**
  - UDP header is only eight-bytes long

# Popular Applications That Use UDP

- **Multimedia streaming**
  - Retransmitting packets is not always worthwhile
  - E.g., phone calls, video conferencing, gaming, IPTV
- **Simple query-response protocols**
  - Overhead of connection establishment is overkill
  - E.g., Domain Name System (DNS), DHCP, etc.



How to introduce reliability in transport?

What are the expectations from a reliable packet delivery system?

- Guaranteed delivery
- Ordered delivery
- At-most one copy of the message
- (any or all of the above)



What are the features of the underlying 'best effort' network?

- Packet loss
- Out of order delivery
- Multiple copies
- Delay variations
- A non-zero probability of packet delivery

# Transmission Control Protocol (TCP)

# TCP “Segments”


Host A



Segment sent when:

1. Segment full (Max Segment Size),
2. Not full, but times out, or
3. “Pushed” by application.


Host B



# Sequence Number

Host A

ISN (initial sequence number)



Sequence  
number =  
 $\text{ISN} + \# \text{ of the } 1^{\text{st}} \text{ byte the segment}$

TCP Data

TCP Data

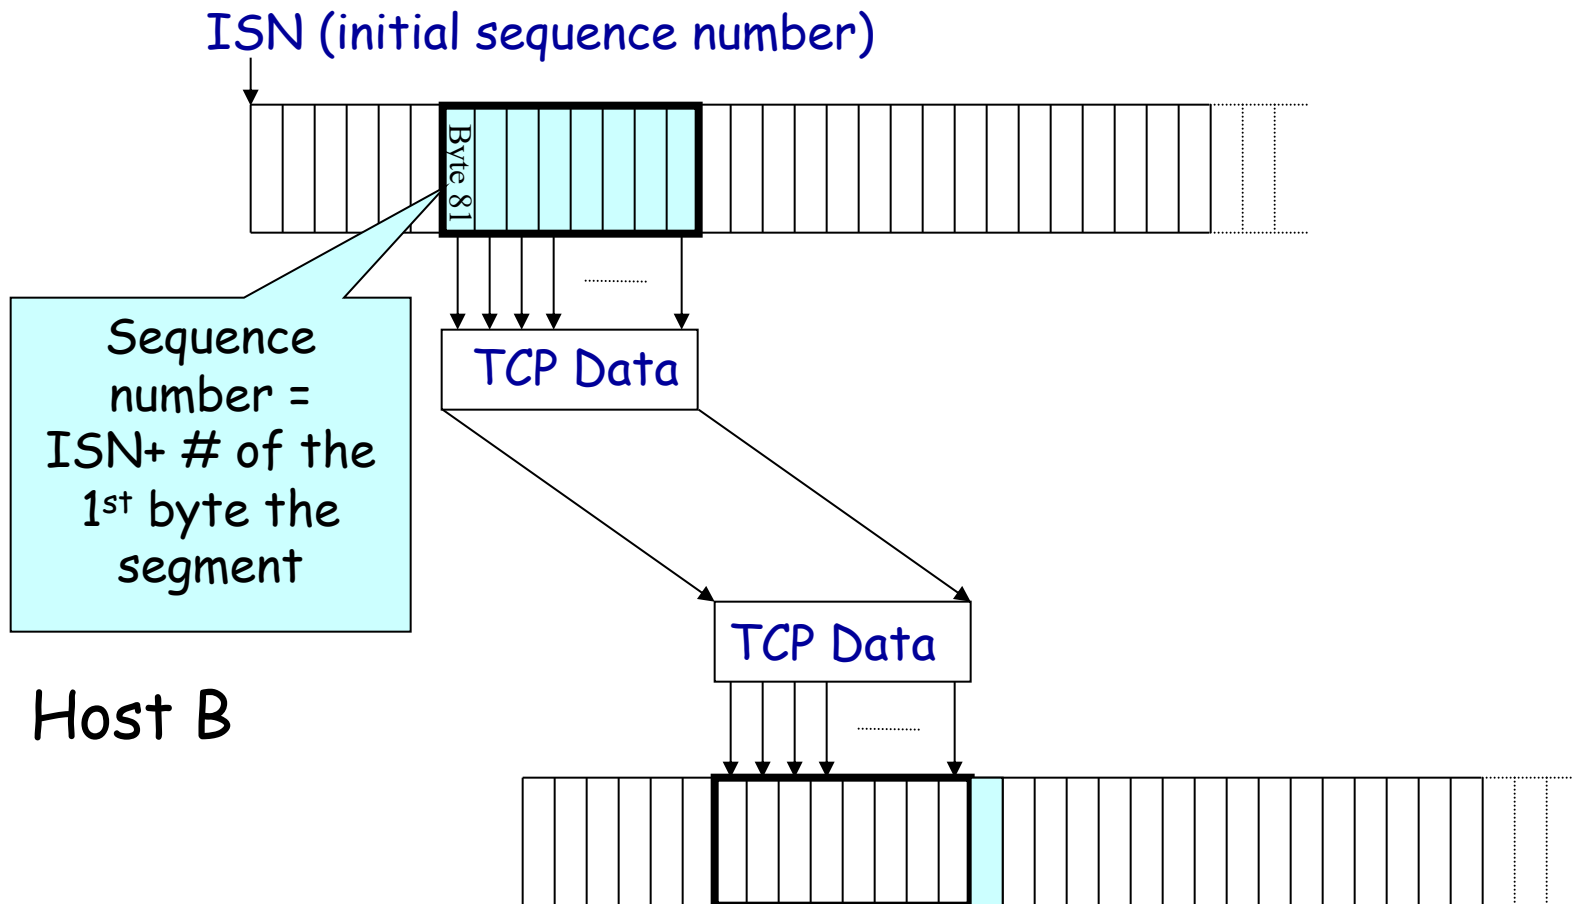
Host B





# Sequence Number

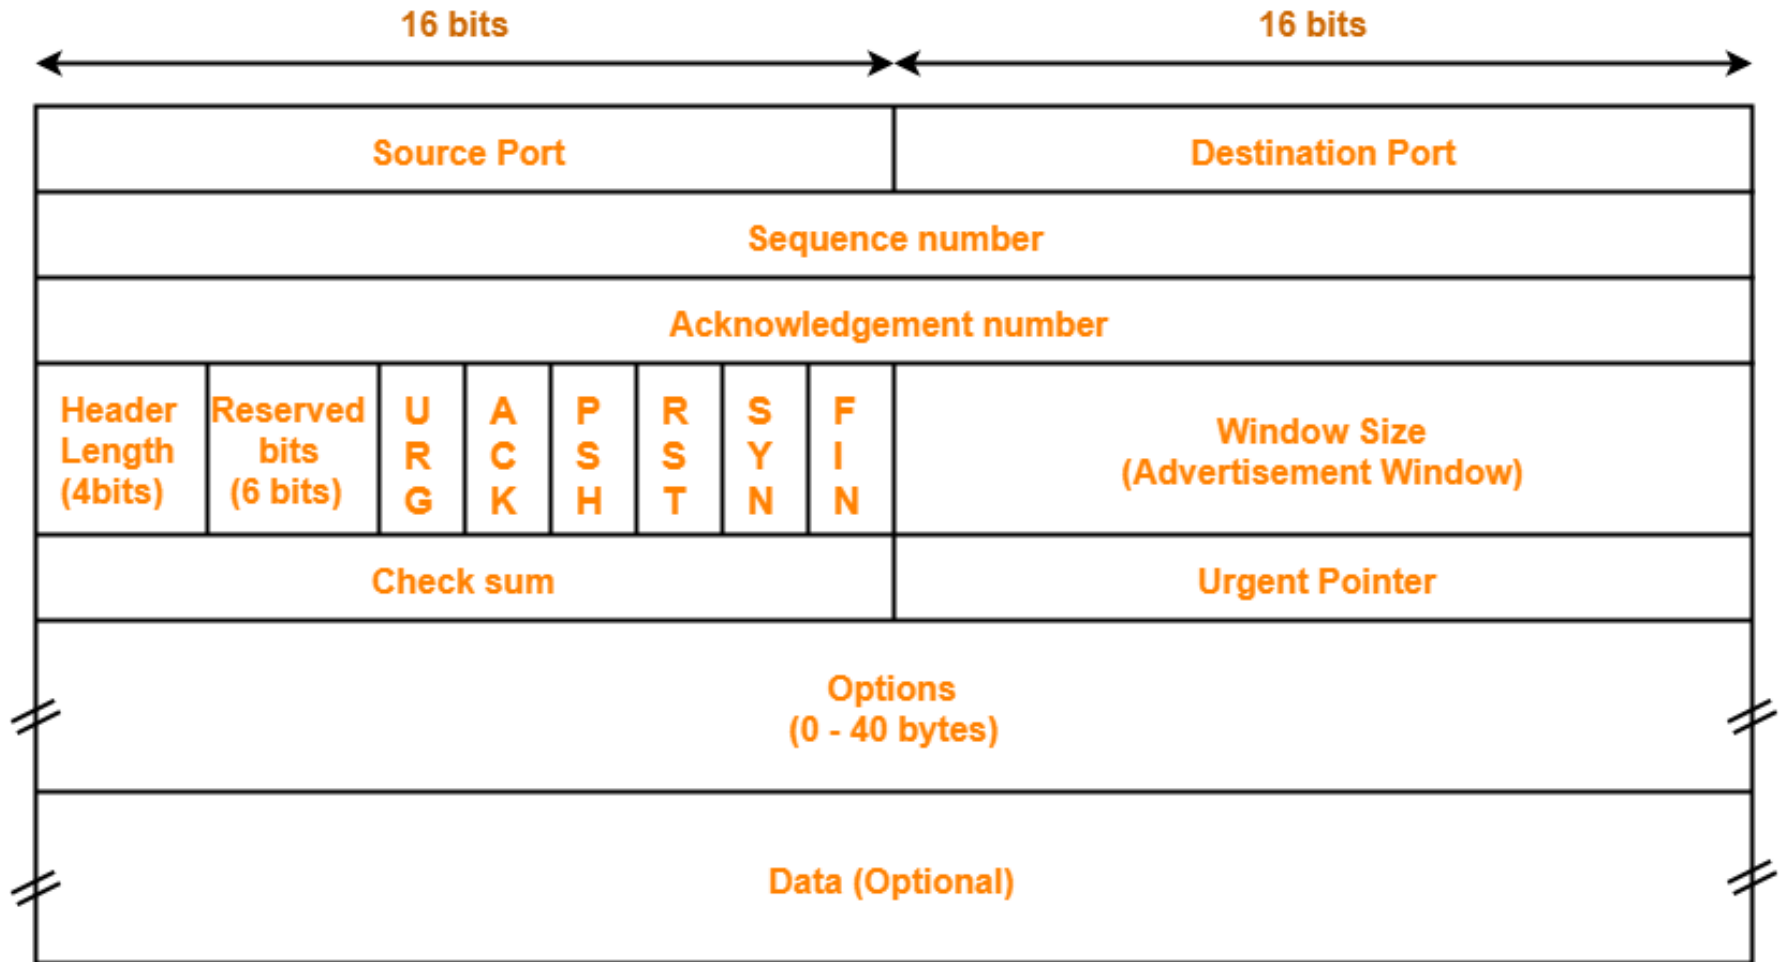
Host A



Host B

How are acknowledgements identified?

# TCP Header

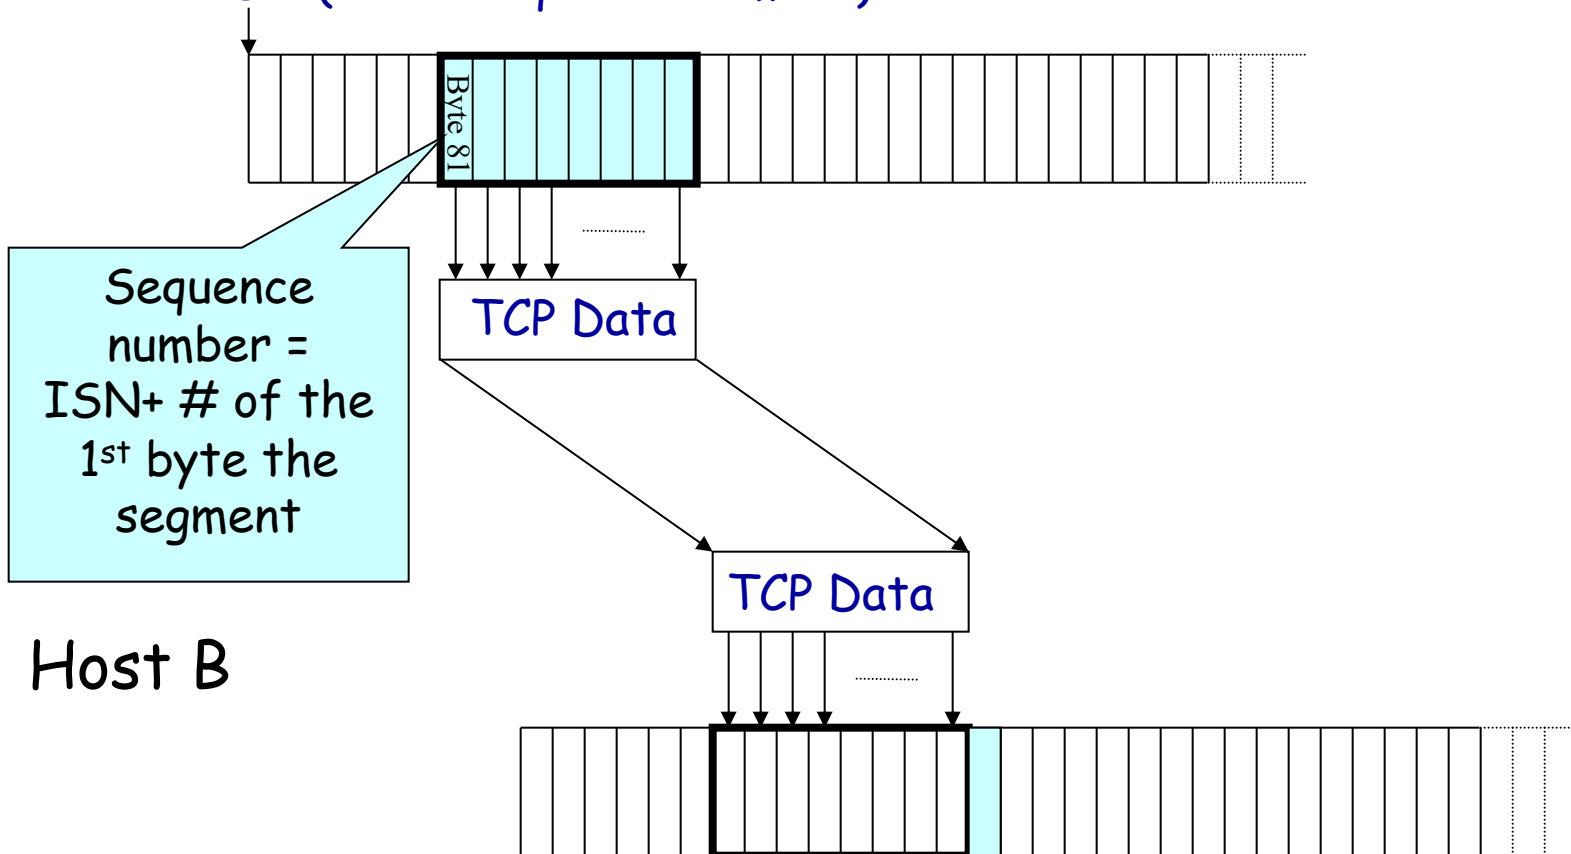


**TCP Header**

# Acknowledgement Number

Host A


ISN (initial sequence number)



Host B

# Starting and Ending a Connection: TCP Handshakes (Three-Way Handshake)

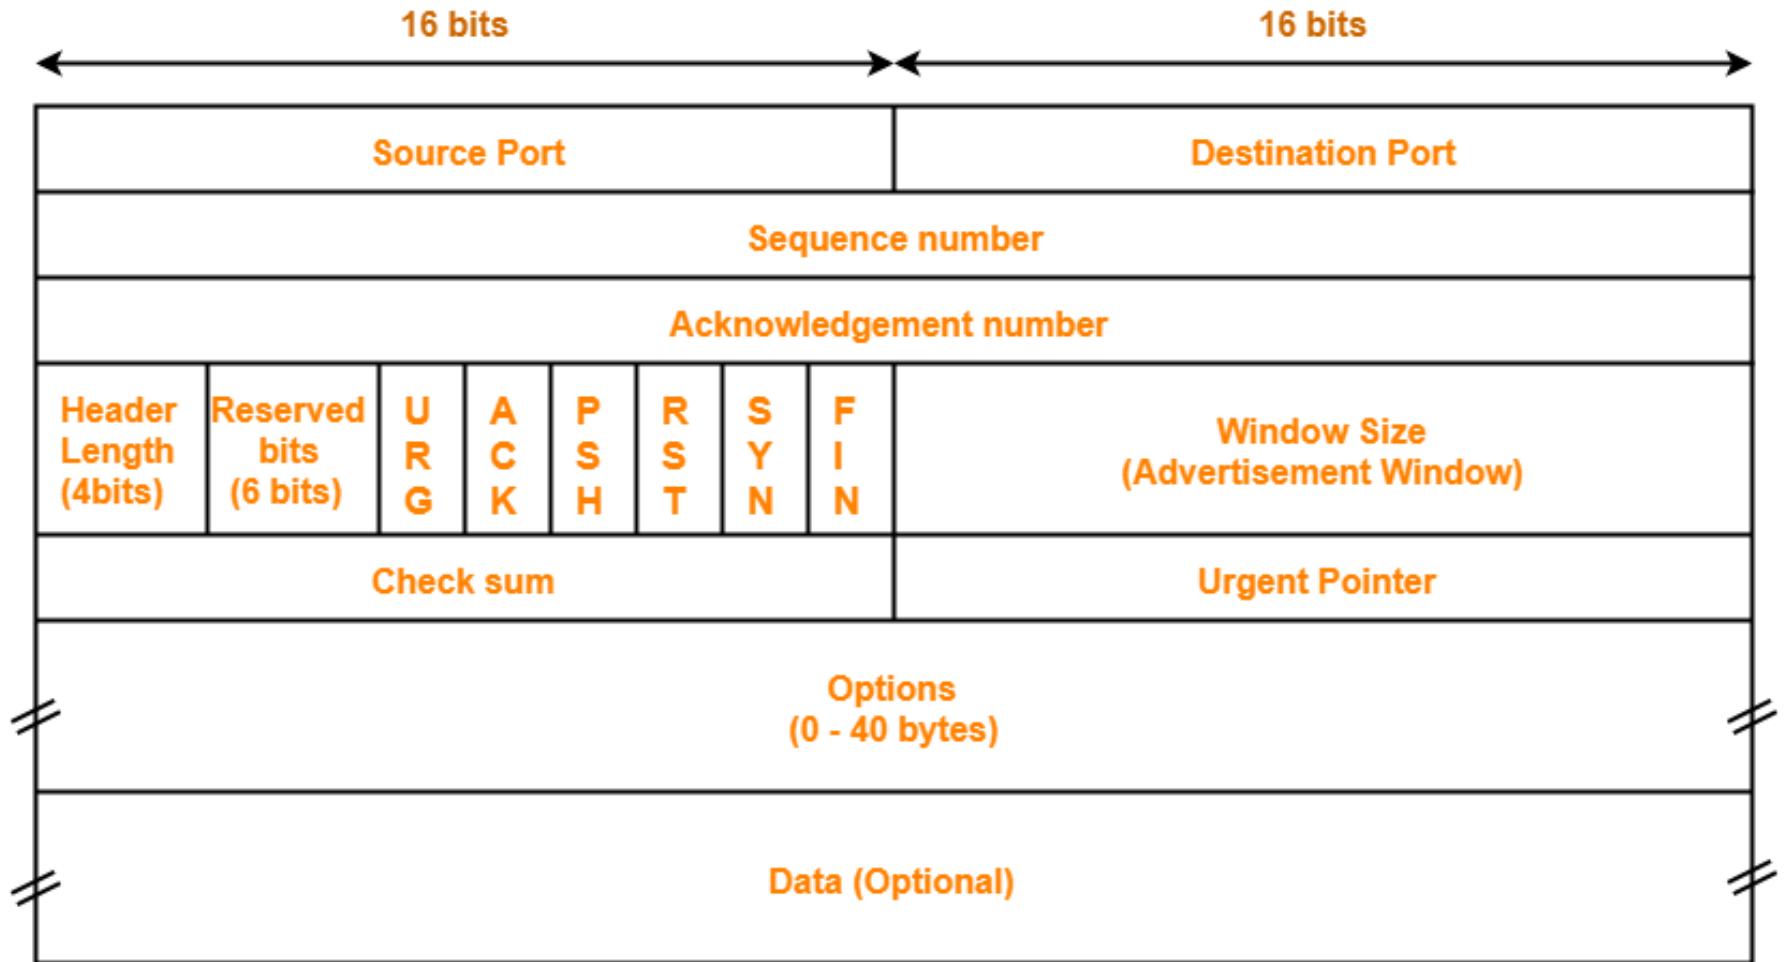
# Establishing a TCP Connection




Each host tells  
its ISN to the  
other host.

- Three-way handshake to establish connection
  - Host A sends a **SYN** (open) to the host B
  - Host B returns a SYN acknowledgment (**SYN ACK**)
  - Host A sends an **ACK** to acknowledge the SYN ACK

# TCP Header



**TCP Header**



RISN - Random initial sequence number  
SYN flag consumes 1 sequence number  
ACK flag consumes 0 sequence number  
Databyte consumes 1 sequence number

**SENDER**

**RECEIVER**

RISN = 521

①

Seq. = 521 , SYN = 1 , MSS = 1460B  
window = 14600B

②

RISN

Seq. = 2000 , SYN = 1 , MSS = 500B  
window = 10000B , Ack no. = 522 ,  
ACK = 1

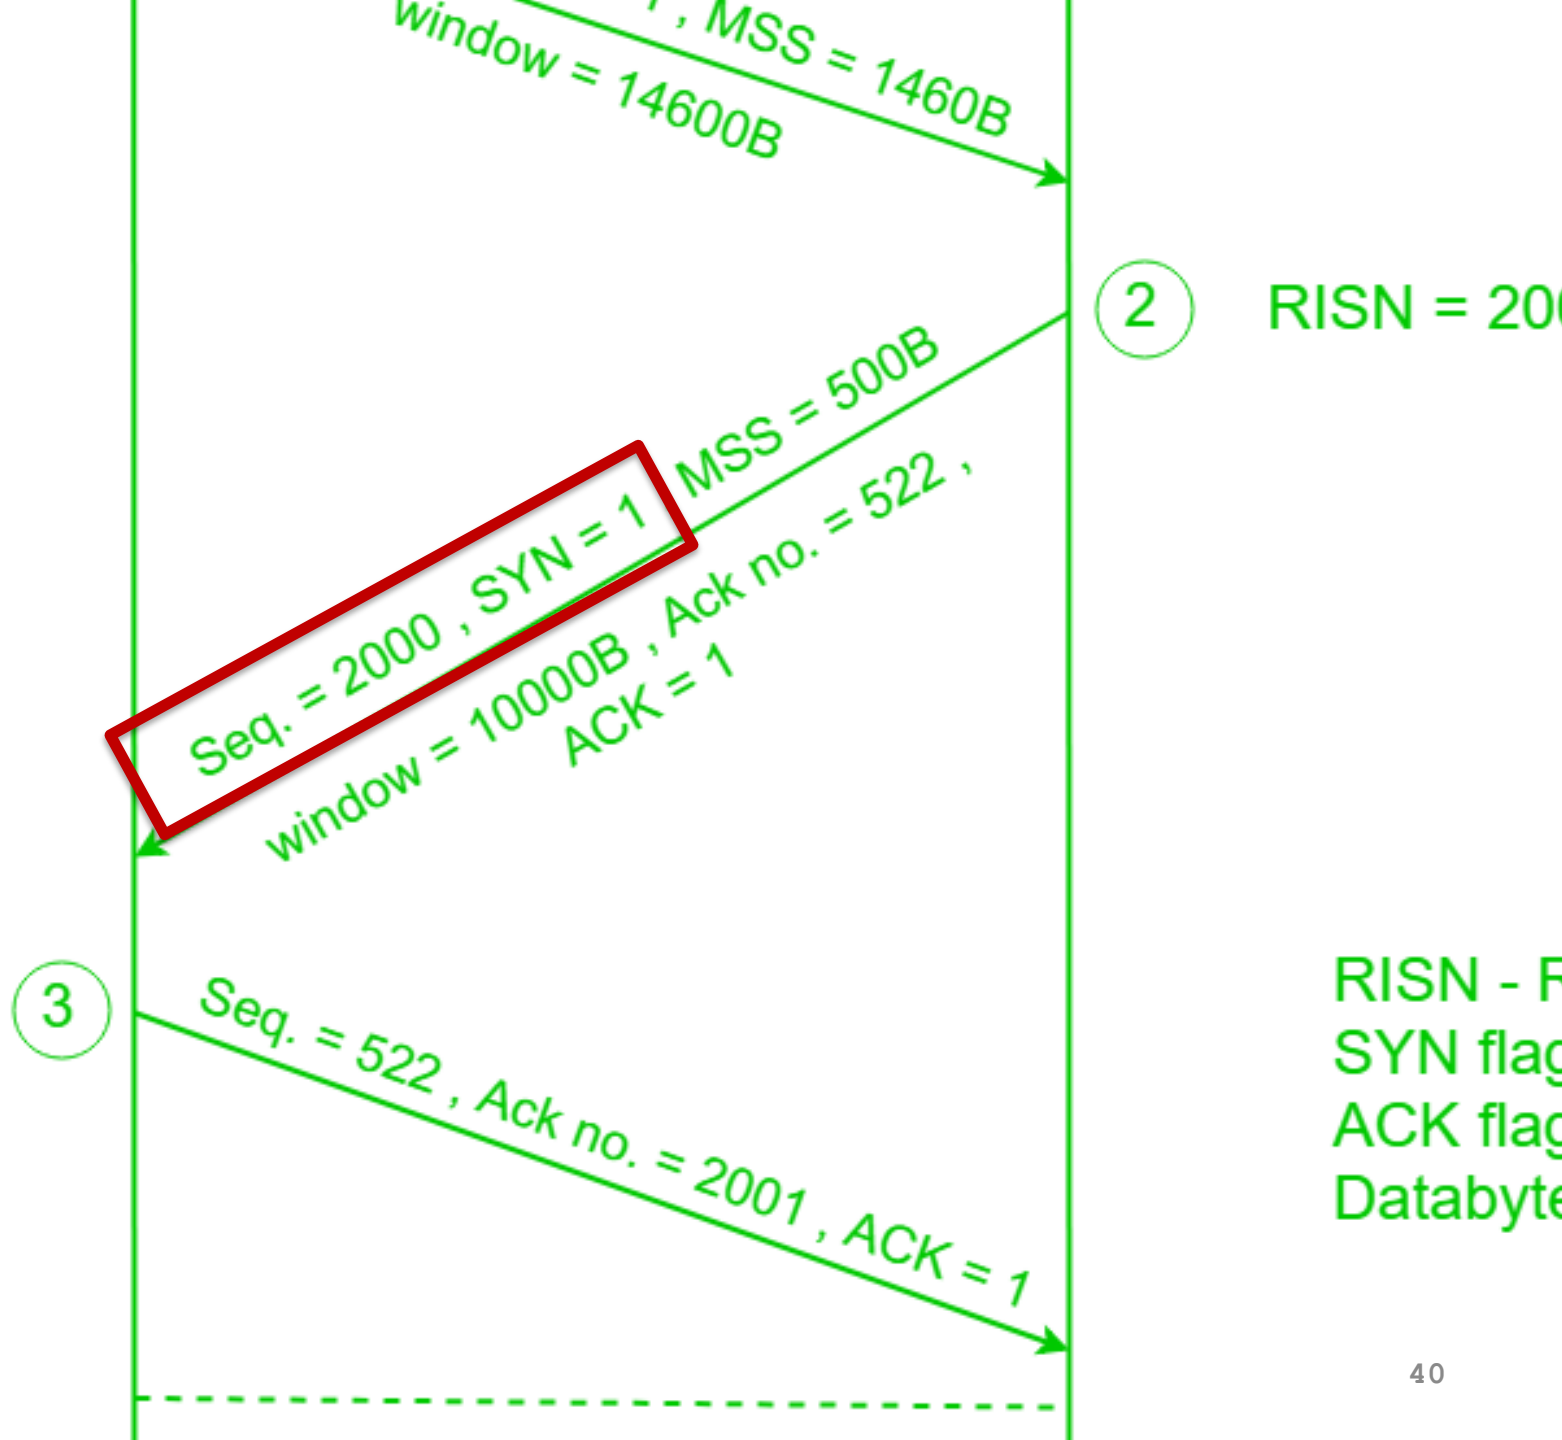


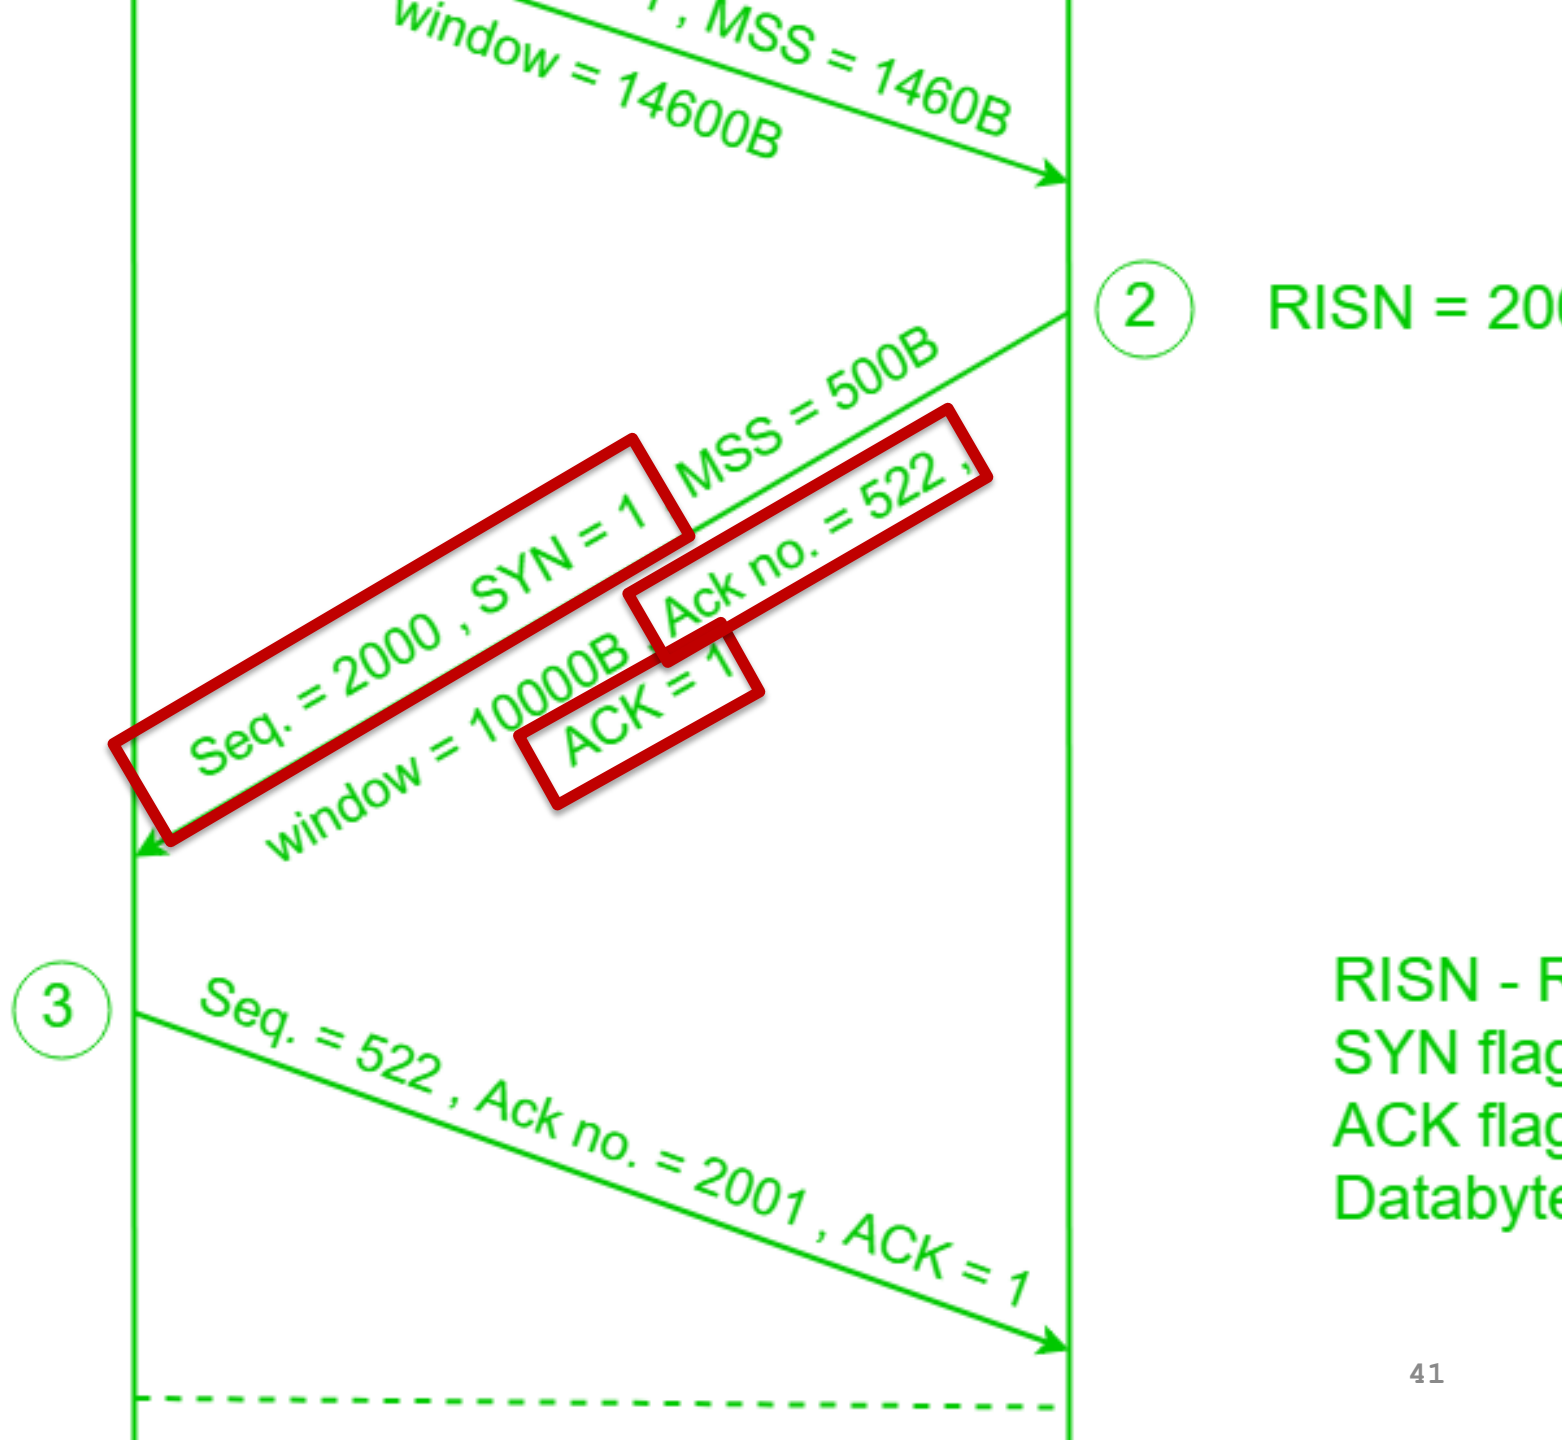
# Step 1: A' s Initial SYN Packet

Flags: **SYN**  
FIN  
RST  
PSH  
URG  
ACK

A' s port		B' s port	
A' s Initial Sequence Number			
Acknowledgment			
20	0	Flags	Advertised window
Checksum			Urgent pointer
Options (variable)			

**A tells B it wants to open a connection...**






# Step 2: B' s SYN-ACK Packet

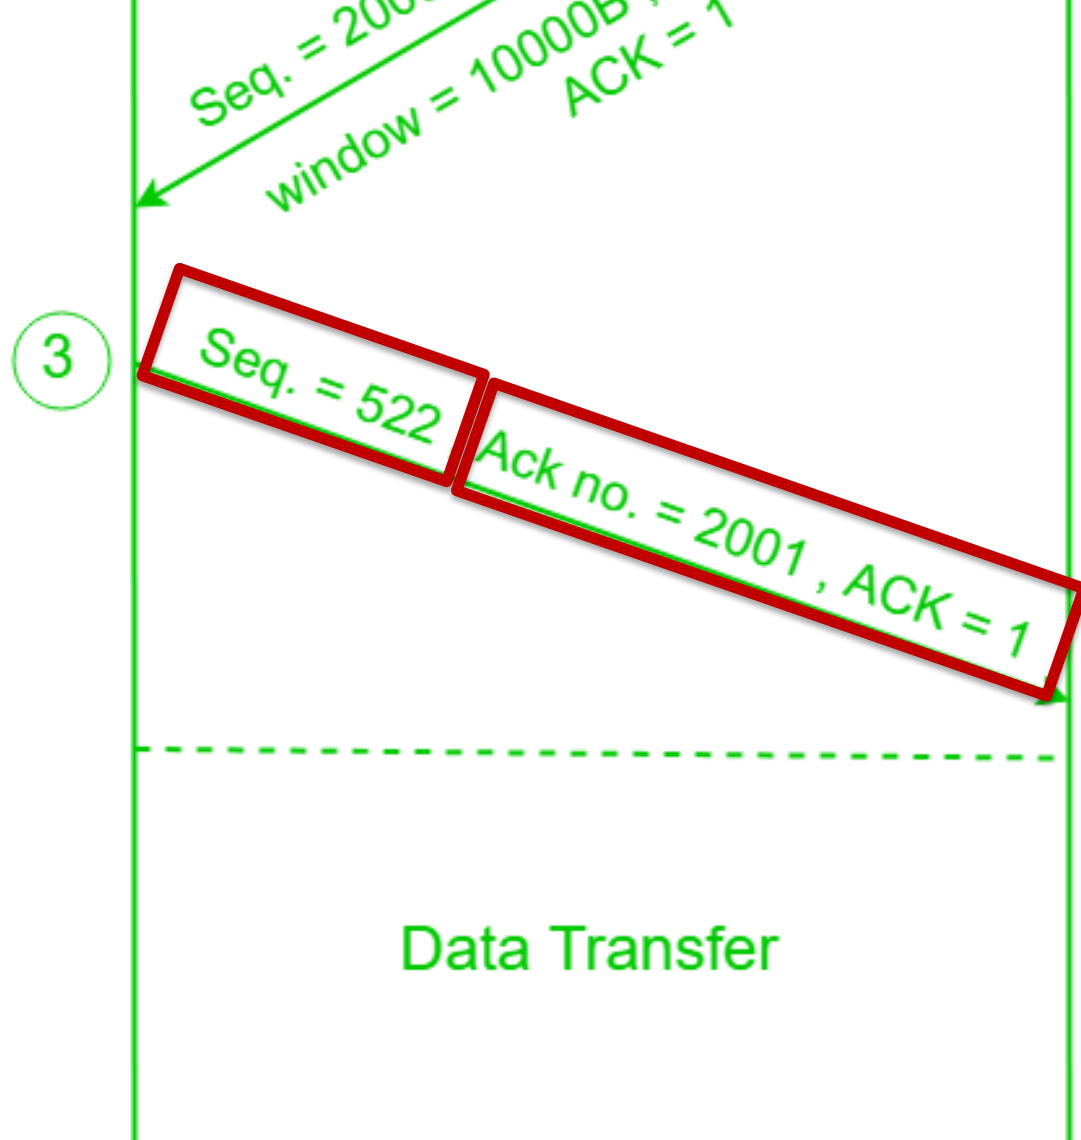
Flags: **SYN**  
FIN  
RST  
PSH  
URG  
**ACK**

B' s port		A' s port	
B' s Initial Sequence Number			
A' s ISN plus 1			
20	0	Flags	Advertised window
Checksum			Urgent pointer
Options (variable)			

**B tells A it accepts, and is ready to hear the next byte...  
... upon receiving this packet, A can start sending data**



RISN - R  
SYN flag  
ACK flag  
Databyte



RISN - R  
SYN flag  
ACK flag  
Databyte

# Step 3: A' s ACK of the SYN-ACK

Flags: SYN  
FIN  
RST  
PSH  
URG  
**ACK**

A' s port		B' s port	
Sequence number			
B' s ISN plus 1			
20	0	Flags	Advertised window
Checksum			Urgent pointer
Options (variable)			

**A tells B it is okay to start sending**  
**... upon receiving this packet, B can start sending data**

**SENDER**

**RECEIVER**

RISN = 521

①

Seq. = 521 , SYN = 1 , MSS = 1460B  
window = 14600B

②

RISN = 2000

Seq. = 2000 , SYN = 1 , MSS = 500B  
window = 10000B , Ack no. = 522 ,  
ACK = 1

③

Seq. = 522 , Ack no. = 2001 , ACK = 1

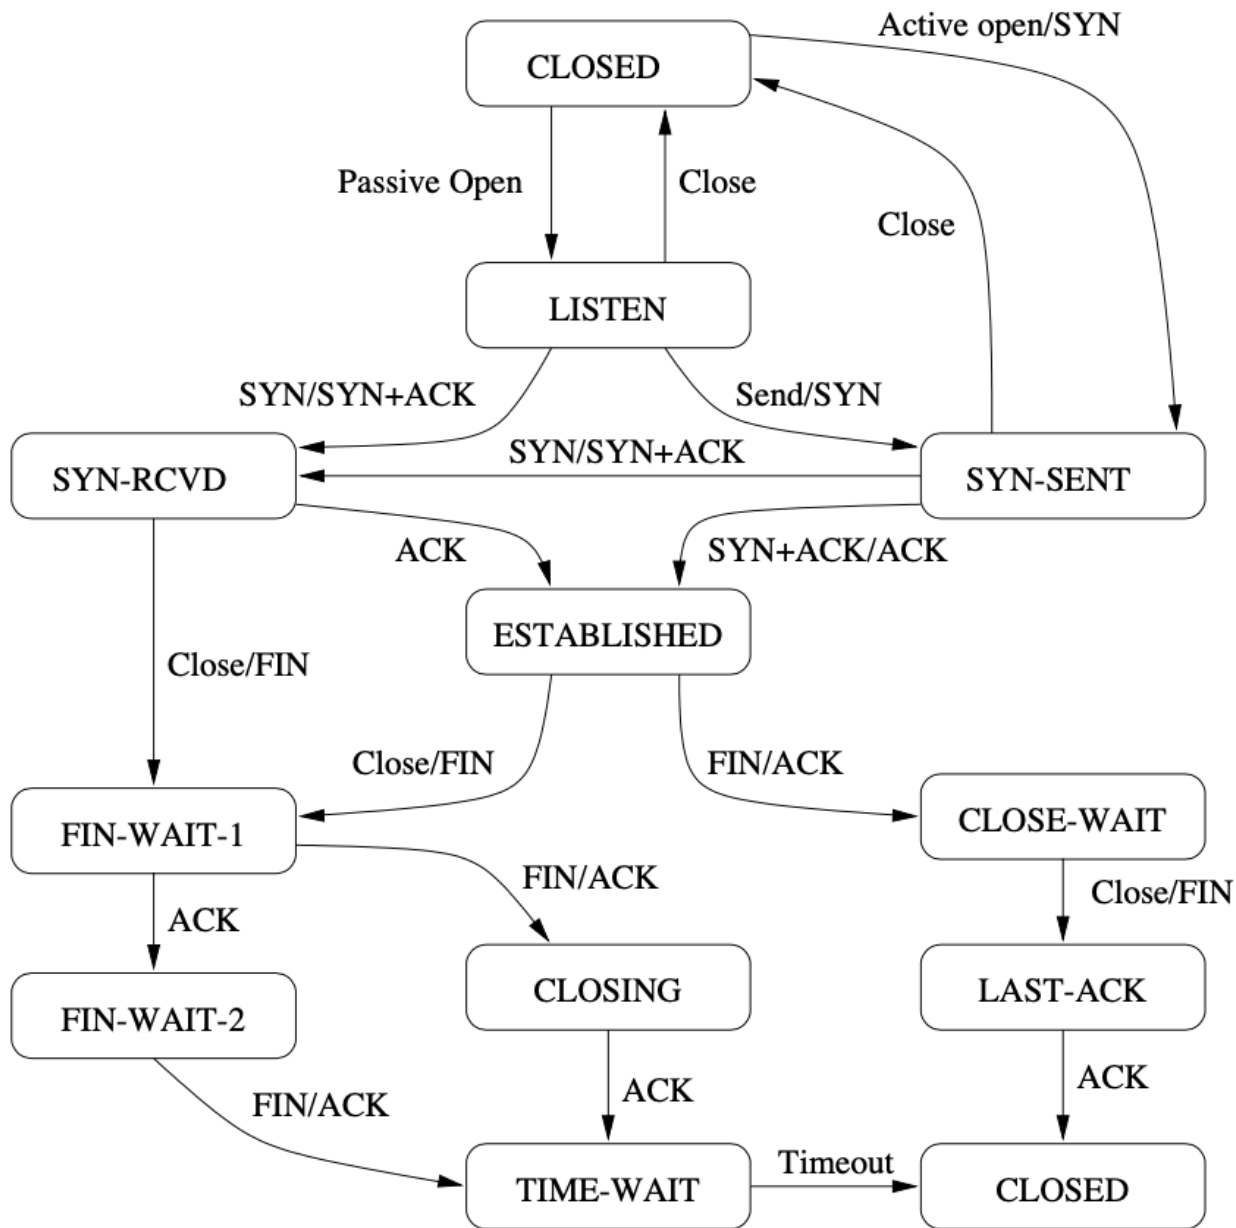
RISN - Random initial sequence number  
SYN flag consumes 1 sequence number  
ACK flag consumes 0 sequence number  
Databyte consumes 1 sequence number

Data Transfer



# What is really happening?

- Exchange of ISNs
  - (a) Alice --> Bob SYNchronize with my Initial Sequence Number of X
  - (b) Alice <-- Bob I received your syn, I ACKnowledge that I am ready for [X+1]
  - (c) Alice <-- Bob SYNchronize with my Initial Sequence Number of Y
  - (d) Alice --> Bob I received your syn, I ACKnowledge that I am ready for [Y+1]
- Piggybacking information
  - (a) Alice --> Bob : SYN
  - (b) Alice <-- Bob : SYN+ACK
  - (c) Alice --> Bob : ACK
- Can data piggyback on 3-way handshake packets?




**Fig. 4.** TCP state-transition diagram

# What if the SYN Packet Gets Lost?


- Suppose the SYN packet gets lost
  - Packet is lost inside the network, or
  - Server rejects the packet (e.g., listen queue is full)
- Eventually, no SYN-ACK arrives
  - Sender sets a timer and wait for the SYN-ACK
  - ... and retransmits the SYN if needed
- How should the TCP sender set the timer?
  - Sender has no idea how far away the receiver is
  - Some TCPs use a default of 3 or 6 seconds

# Tearing Down the Connection



- Closing (each end of) the connection
  - Finish (FIN) to close and receive remaining bytes
  - And other host sends a FIN ACK to acknowledge

# Tearing Down the Connection

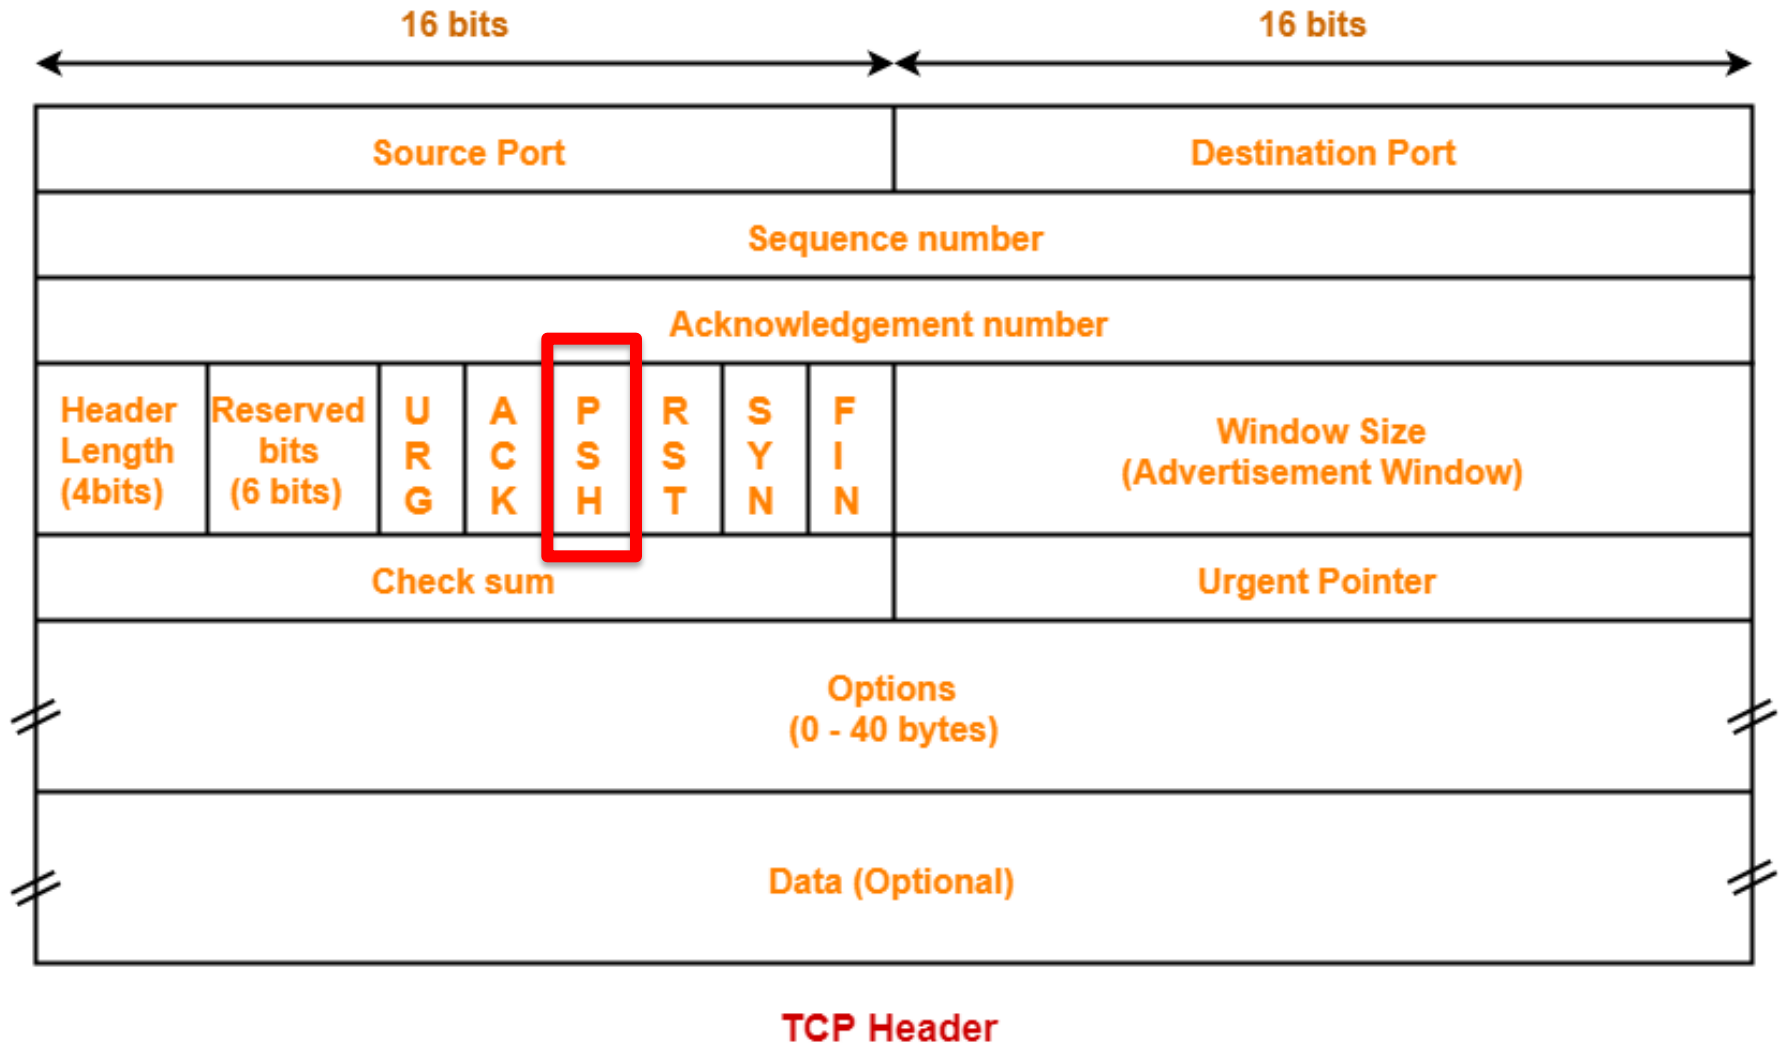


- Closing (each end of) the connection
  - Finish (FIN) to close and receive remaining bytes
  - And other host sends a FIN ACK to acknowledge
  - Reset (RST) to close and not receive remaining bytes

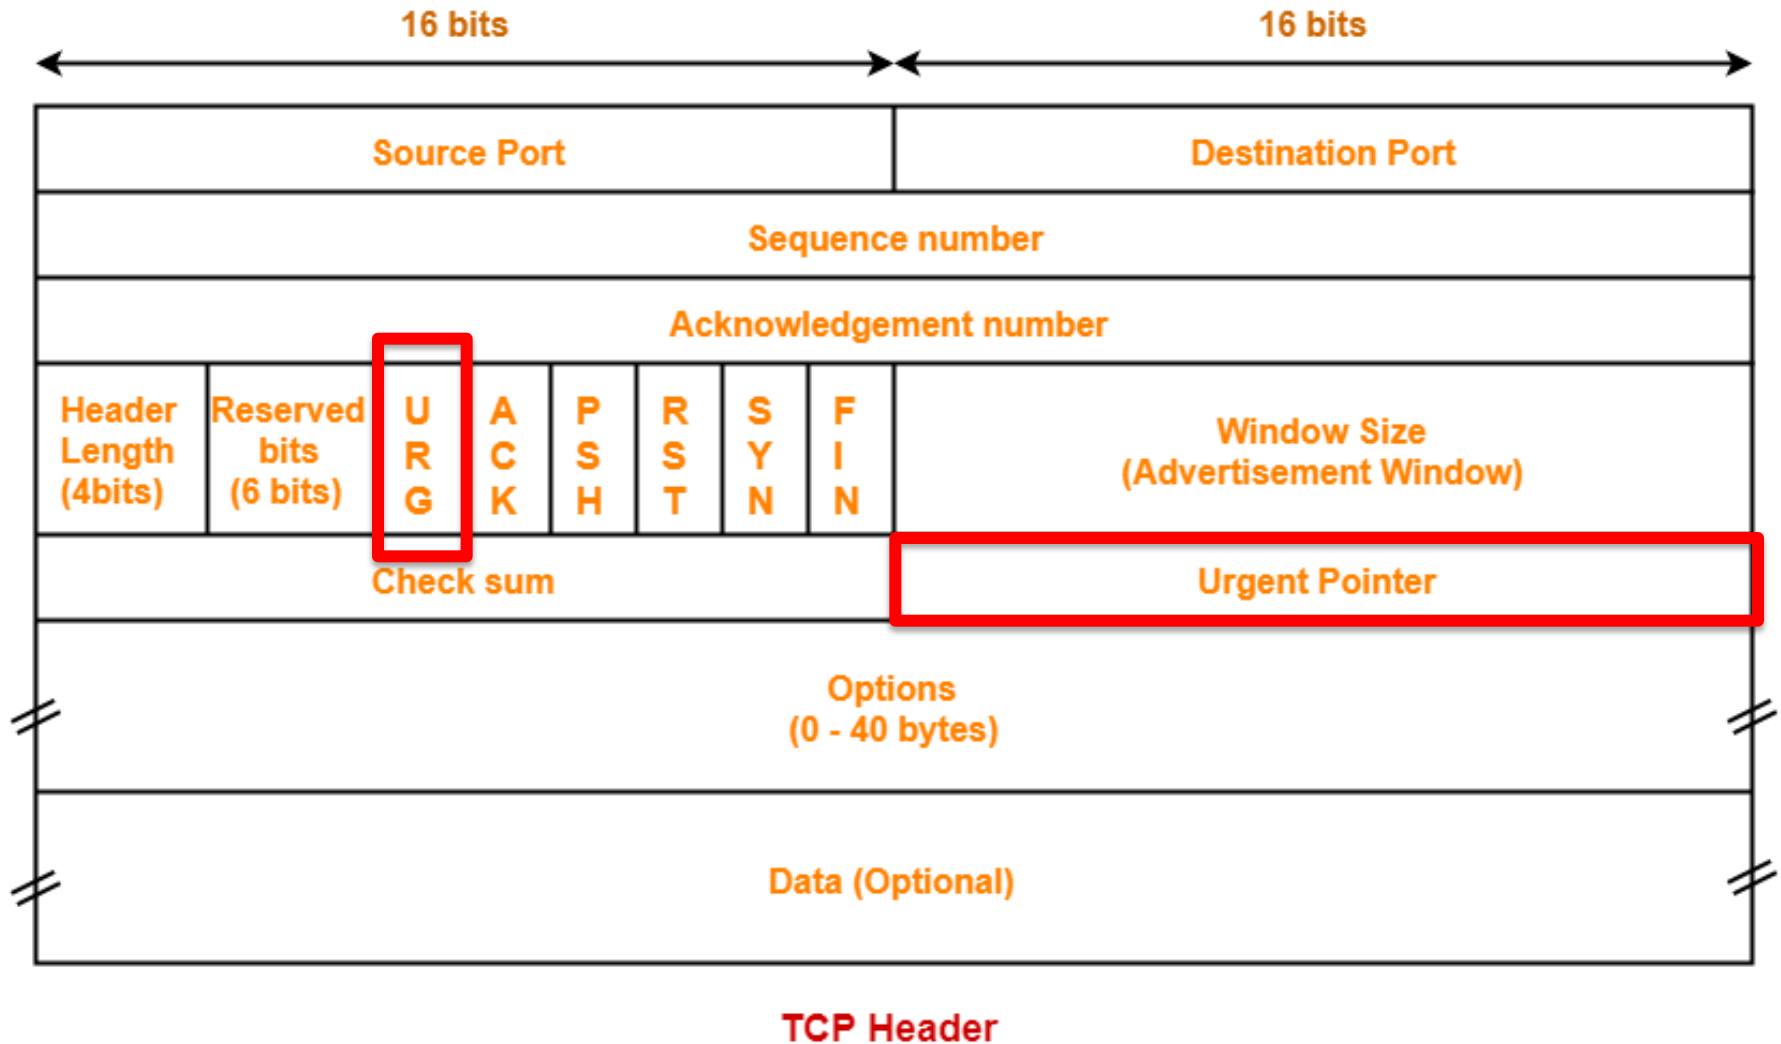
# Sending/Receiving the FIN Packet

- **Sending a FIN: close()**
  - Process is done sending data via the socket
  - Process invokes “close()” to close the socket
  - Once TCP has sent all the outstanding bytes...
  - ... then TCP sends a FIN
- **Receiving a FIN: EOF**
  - Process is reading data from the socket
  - Eventually, the attempt to read returns an EOF

# TCP Header: PSH and URG flags



# TCP Header: PSH and URG flags





# Reliable Delivery on a Lossy Channel With Bit Errors

# Challenges of Reliable Data Transfer

- Over a perfectly reliable channel
  - Easy: sender sends, and receiver receives
- Over a channel with bit errors
  - Receiver detects errors and requests retransmission
- Over a lossy channel with bit errors
  - Some data are missing, and others corrupted
  - Receiver cannot always detect loss
- Over a channel that may reorder packets
  - Receiver cannot distinguish loss from out-of-order