



Computer Networks

CMSC 417 : Spring 2024



COMPUTER SCIENCE
UNIVERSITY OF MARYLAND

Topic: Transport Layer Protocols (TCP)
(Textbook chapter 5)

Nirupam Roy

Tu-Th 2:00-3:15pm

CSI 2117

March 12th, 2024



Retransmission hints

- Acknowledgments from receiver
 - Positive: “okay” or “uh huh” or “ACK”
 - Negative: “please repeat that” or “NACK”
- Retransmission by the sender
 - After *not* receiving an “ACK”
 - After receiving a “NACK”
- Timeout by the sender (“stop and wait”)
 - Don’ t wait forever without some acknowledgment


TCP Support for Reliable Delivery

- **Detect bit errors: checksum**
 - Used to detect corrupted data at the receiver
 - ...leading the receiver to drop the packet
- **Detect missing data: sequence number**
 - Used to detect a gap in the stream of bytes
 - ... and for putting the data back in order
- **Recover from lost data: retransmission**
 - Sender retransmits lost or corrupted data
 - Two main ways to detect lost packets

TCP Acknowledgments

Host A

ISN (initial sequence number)




Sequence number
= ISN+1st byte

TCP Data

ACK sequence
number = next
expected byte


TCP Data

Host B



Automatic Repeat reQuest (ARQ)

- **ACK and timeouts**
 - Receiver sends ACK when it receives packet
 - Sender waits for ACK and times out
- **Simplest ARQ protocol**
 - Stop and wait
 - Send a packet, stop and wait until ACK arrives

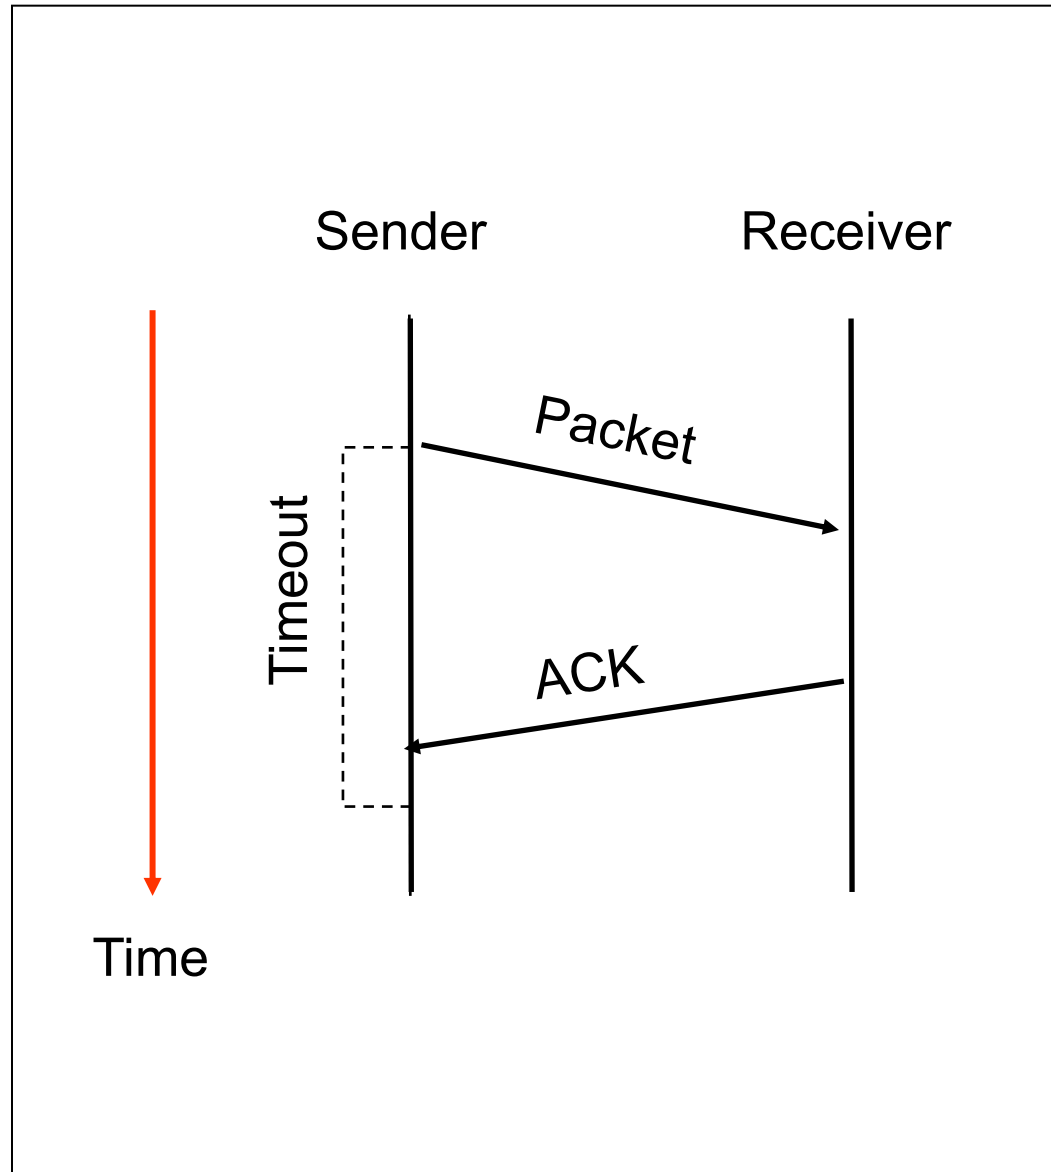


Automatic Repeat reQuest (ARQ)

- We will discuss two variations
 1. Stop and Wait
 2. Sliding window


Reliable Delivery on a Lossy Channel With Bit Errors

Stop-and-Wait Protocol



- A few design decisions
 1. Unique packet identifiers: Why and how
 2. Duplicate packets: Why and how to eliminate
 3. Timeout: picking the ideal value

- A few design decisions
 1. Unique packet identifiers: Why and how
 2. Duplicate packets: Why and how to eliminate
 3. Timeout: picking the ideal value



- Wrap around is possible

Network	B*8 bits/sec	B bytes/sec	Twrap secs
ARPA NET	56 kbps	7 KBps	3*10**5 (~3.6 days)
DS1	1.5 Mbps	190 KBps	10**4 (~3 hours)
Ethernet	10 Mbps	1.25 MBps	1700 (~30 mins)
DS3	45 Mbps	5.6 MBps	380
FDDI	100 Mbps	12.5 MBps	170
Gigabit	1 Gbps	125 MBps	17

- Wrap around-aware sequence number comparison

If A and B are sequence numbers,

$A < B$ if $0 < (B - A) < 2^{31}$,

computed in unsigned 32-bit arithmetic

- PAWS: Protect Against Wrapped Sequence Numbers (RFC 1323) – Use sequence number and timestamps.

- A few design decisions
 1. Unique packet identifiers: Why and how
 2. Duplicate packets: Why and how to eliminate
 3. Timeout: picking the ideal value

- A few design decisions

2. Duplicate packets: Why and how to eliminate

- Receiver maintain the last in-sequence byte number in a variable, say `rcv_seqnum`
- If the Receiver receives a packet with seq number `less than or equal to rcv_seqnum`

- A few design decisions

2. Duplicate packets: Why and how to eliminate

- Receiver maintain the last in-sequence byte number in a variable, say `rcv_seqnum`
- If the Receiver receives a packet with seq number **less than or equal to `rcv_seqnum`**
 - >>It sends an ACK for that packet, and discards the packet
- If the Receiver receives a packet with seq number **(`rcv_seqnum+1`):**


- A few design decisions

2. Duplicate packets: Why and how to eliminate


- Receiver maintain the last in-sequence byte number in a variable, say `rcv_seqnum`
- If the Receiver receives a packet with seq number **less than or equal to `rcv_seqnum`**
 - >> It sends an ACK for that packet, and discards the packet
- If the Receiver receives a packet with seq number **(`rcv_seqnum+1`):**
 - >> It sends an ACK for that packet, and includes the packet to the application buffer
- What if the receiver gets a packet with seq number **greater than (`rcv_seqnum+1`) ?**

- A few design decisions
 1. Unique packet identifiers: Why and how
 2. Duplicate packets: Why and how to eliminate
 3. Timeout: picking the ideal value


Reasons for Retransmission



Packet lost



ACK lost
DUPLICATE
PACKET



Early timeout
DUPLICATE
PACKETS

How Long Should Sender Wait?

- Sender sets a timeout to wait for an ACK
 - Too short:
 - >> wasted retransmissions
 - Too long:
 - >> excessive delays when packet lost


How Long Should Sender Wait?

- Sender sets a timeout to wait for an ACK
 - Too short:
 - >> wasted retransmissions
 - Too long:
 - >> excessive delays when packet lost
- TCP sets timeout as a function of the RTT
 - Expect ACK to arrive after a “round-trip time”
 - ... plus a fudge factor to account for queuing
- But, how does the sender know the RTT?

How Long Should Sender Wait?


- Sender sets a timeout to wait for an ACK
 - Too short:
 - >> wasted retransmissions
 - Too long:
 - >> excessive delays when packet lost
- TCP sets timeout as a function of the RTT
 - Expect ACK to arrive after a “round-trip time”
 - ... plus a fudge factor to account for queuing
- But, how does the sender know the RTT?
 - Running average of delay to receive an ACK

TCP timestamp




TCP timestamp

Sender's current time



TCP Timestamp Option (10 bytes)


Sender's current time



Copied timestamp from the packet
being acknowledged.

(valid only when the packet contains an ACK.
i.e., the ACK bit is set in TCP header).

Example RTT Estimation




Reliable Delivery on a Lossy Channel With Bit Errors


TCP Flow control: Sliding Window Protocol

Motivation for Sliding Window


- Stop-and-wait is inefficient
 - Only one TCP segment is “in flight” at a time
 - Especially bad for high “delay-bandwidth product”




Revisiting “delay X bandwidth”




Revisiting “delay X bandwidth”



Revisiting “delay X bandwidth”




Revisiting “delay X bandwidth”




Consider two-way delay
or round-trip delay

Revisiting “delay X bandwidth”



stop-and-wait operation



$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Packet size = L bits

Network bandwidth = R bits-per-sec

Performance – an implementation of stop-and-wait

- ❑ Stop-and-wait works, but performance stinks
- ❑ ex: 1 Gbps link, 15 ms prop. delay, 8kbit (8000 bit) packet:


$$d_{trans} = \frac{L}{R} = \frac{8000\text{bits}}{10^9\text{bps}} = 8\text{microseconds}$$

- U_{sender} : **utilization** - fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec -> 33kB/sec thruput over 1 Gbps link
- network protocol limits use of physical resources!

stop-and-wait operation




$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Packet size = L bits

Network bandwidth = R bits-per-sec

stop-and-wait operation




$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

Packet size = L bits

Network bandwidth = R bits-per-sec

stop-and-wait operation



Network protocol limits the full utilization of physical resources.

Transport-layer vs Link-layer capacity gap


Network bandwidth = R bits-per-sec

Transport-layer capacity = 0.00


Pipelined protocols

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver



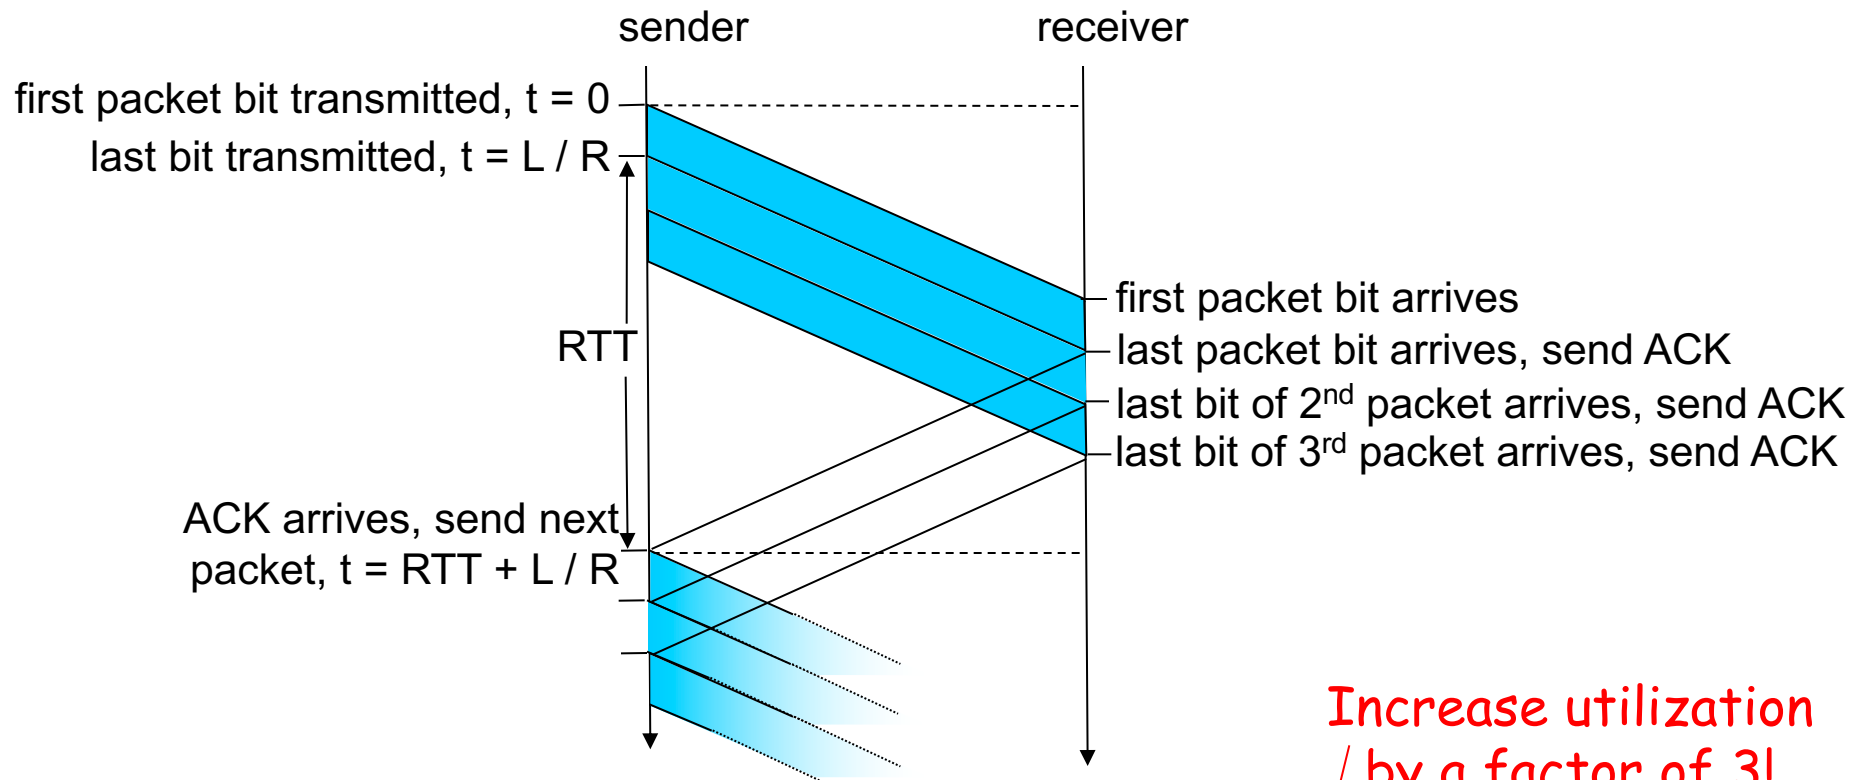
(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation

- Two generic forms of pipelined protocols (or sliding window protocol) depending on the retransmission strategy: *go-Back-N, selective repeat*


Pipelining: increased utilization




Increase utilization
by a factor of 3!

$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$


Revisiting “delay X bandwidth”



Stop-and-wait




Sliding Window




Sliding Window Protocol


Sliding Window: A series of packets to be sent reliably



Sliding Window: A series of packets to be sent reliably




Sliding Window: A series of packets to be sent reliably



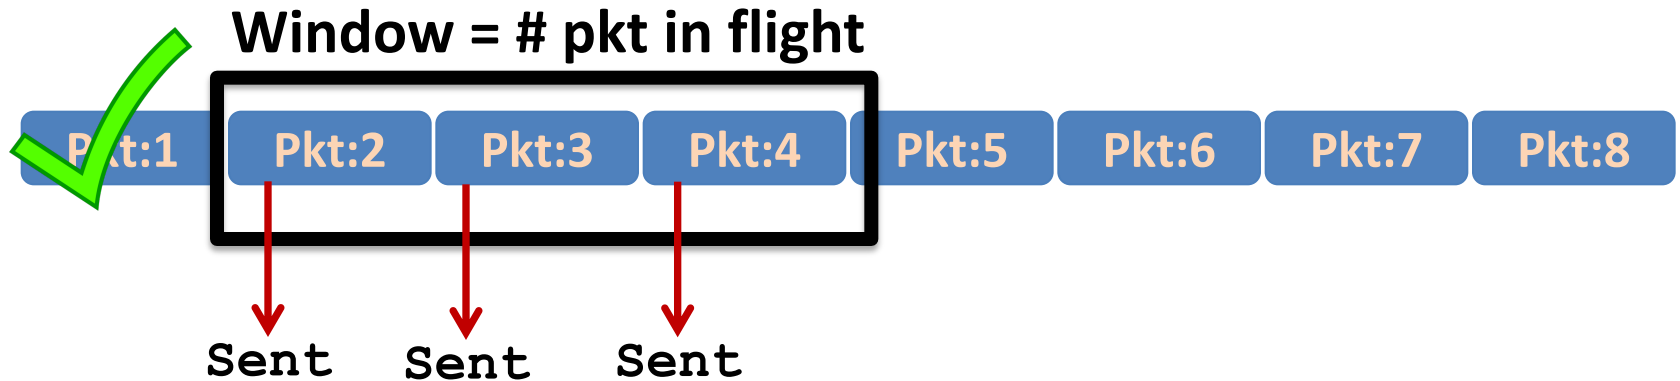
Packet 1 is delivered

Sliding Window: A series of packets to be sent reliably



Packet 1 is delivered. Slide the window.

Sliding Window: A series of packets to be sent reliably



Packet 1 is delivered. Slide the window.

Sliding Window: A series of packets to be sent reliably

