



Computer Networks

CMSC 417 : Spring 2024



COMPUTER SCIENCE
UNIVERSITY OF MARYLAND

**Topic: TCP Congestion-control, TCP history,
Router-assisted congestion control, Link Layer
(Textbook chapter 6 & 2)**

Nirupam Roy


**Tu-Th 2:00-3:15pm
CSI 2117**

March 28th, 2024



TCP Congestion Control

- Additive increase, multiplicative decrease (AIMD)
 - On packet loss, divide congestion window in half
 - On success for last window, increase window linearly



Why Exponential decrease?

- Respond aggressively to bad news
 - Congestion is (very) bad for everyone
 - Need to react aggressively
- Nice theoretical properties
 - Makes efficient use of network resources


TCP Congestion Window

- Each TCP sender maintains a congestion window
 - Max number of bytes to have in transit (not yet ACK' d)
- Adapting the congestion window
 - Decrease upon losing a packet: backing off
 - Increase upon success: optimistically exploring
 - Always struggling to find right transfer rate
- Tradeoff
 - Pro: avoids needing explicit network feedback
 - Con: continually under- and over-shoots “right” rate

Additive Increase, Multiplicative Decrease

- **How much to adapt?**
 - Additive increase: On success of last window of data, increase window by 1 Max Segment Size (MSS)
 - Multiplicative decrease: On loss of packet, divide congestion window in half
- **Much quicker to slow than speed up!**
 - Over-sized windows (causing loss) are much worse than under-sized windows (causing lower throughput)
 - AIMD: A necessary condition for stability of TCP

Leads to the TCP “Sawtooth”




Receiver/Advertised Window vs. Congestion Window

- Flow control
 - Keep a *fast sender* from overwhelming a *slow receiver*
- Congestion control
 - Keep a *set of senders* from overloading the *network*
- Different concepts, but similar mechanisms
 - TCP flow control: receiver window
 - TCP congestion control: congestion window
 - Sender TCP window =
 $\min \{ \text{congestion window, receiver window} \}$

Slow start


How Should a New Flow Start?

Start slow (a small CWND) to avoid overloading network



How Should a New Flow Start?

Start slow (a small CWND) to avoid overloading network




But, could take a long time to get started!

“Slow Start” Phase


- Start with a small congestion window
 - Initially, CWND is 1 MSS
 - So, initial sending rate is MSS / RTT
- Could be pretty wasteful
 - Might be much less than actual bandwidth
 - Linear increase takes a long time to accelerate
- Slow-start phase (really “fast start”)
 - Sender starts at a slow rate (hence the name)
 - ... but increases rate exponentially until the first loss

Slow Start in Action

Double CWND per round-trip time



Slow Start and the TCP Sawtooth



- Why is this called “slow start”?

TCP originally had *no* congestion control

- Source would start by sending entire receiver window
- Led to congestion collapse!
- “Slow start” is, comparatively, slower

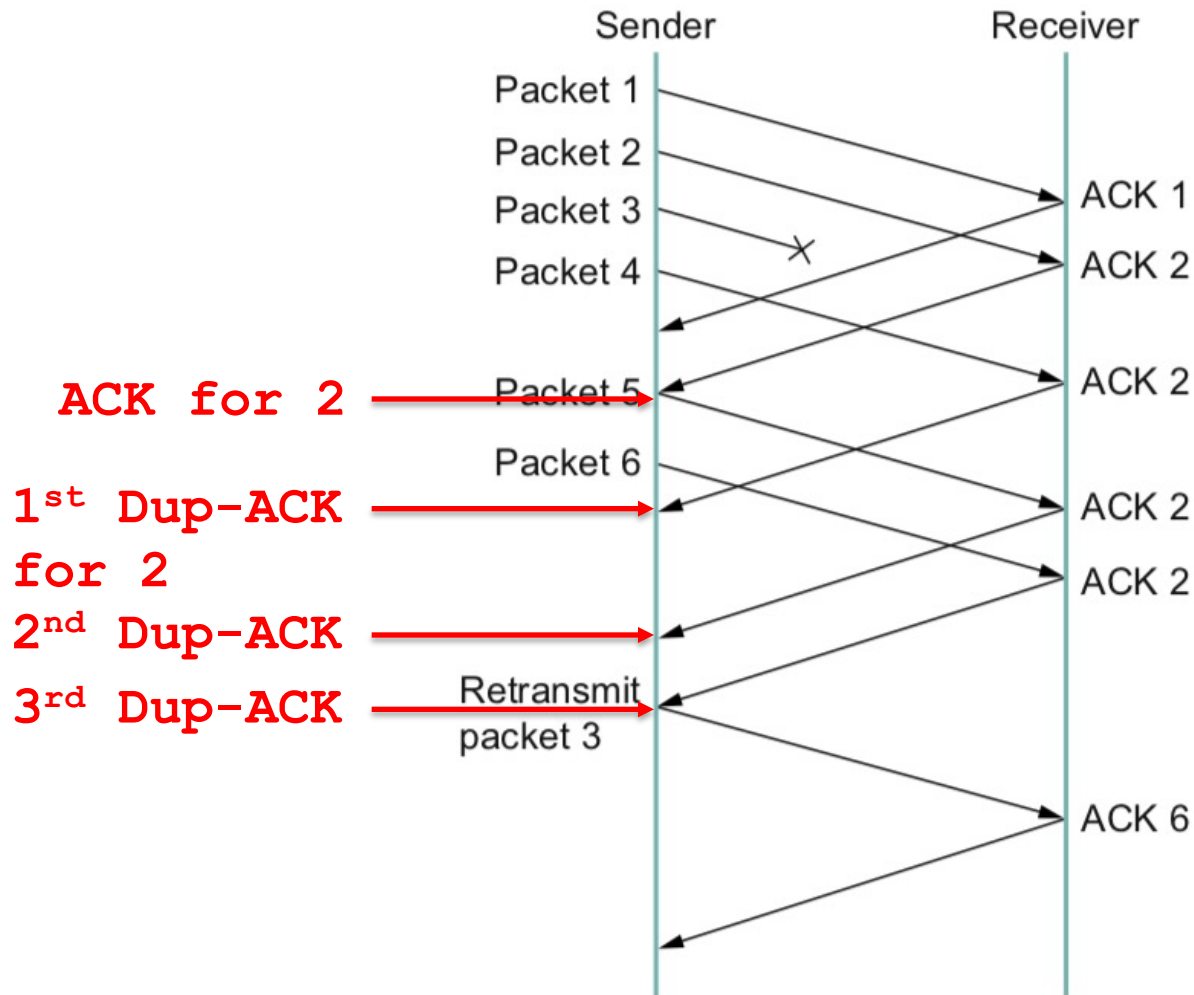
Response to the two types of losses

Response to the two types of losses

(1) Triple duplicate ACK (Fast retransmit)

- Packet n is lost, but packets $n+1$, $n+2$, etc. arrive
- Then, sender quickly resends packet n
- Do a multiplicative decrease and keep going

Fast retransmit

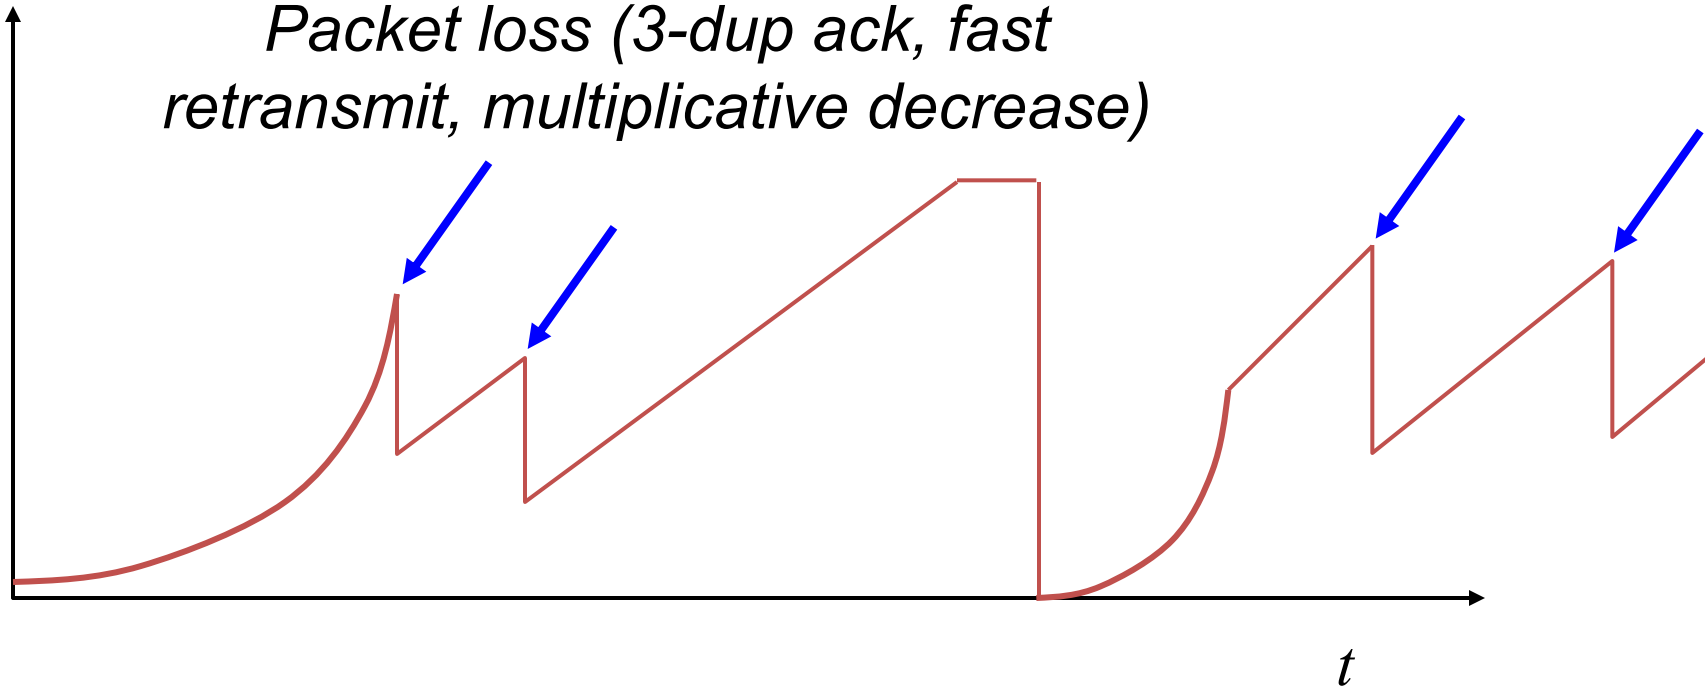


■ FIGURE 6.12 Fast retransmit based on duplicate ACKs.

Congestion window during Fast retransmit

Window

Packet loss (3-dup ack, fast retransmit, multiplicative decrease)



Two Kinds of Loss in TCP


(1) Triple duplicate ACK (Fast retransmit)

- Packet n is lost, but packets $n+1$, $n+2$, etc. arrive
- Then, sender quickly resends packet n
- Do a multiplicative decrease and keep going

(2) Timeout (Repeat slow start – slow start restart)

- Packet n is lost and detected via a timeout
- Blasting entire CWND would cause another burst
- Better to start over with a low CWND

Repeating Slow Start After Timeout (slow-start restart)




Slow-start restart: Go back to CWND of 1, but take advantage of knowing the previous value of CWND.

Repeating Slow Start After Idle Period

- Suppose a TCP connection goes idle for a while
- Eventually, the network conditions change
 - Maybe many more flows are traversing the link
- Dangerous to start transmitting at the old rate
 - Previously-idle TCP sender might blast network
 - ... causing excessive congestion and packet loss
- So, some TCP implementations repeat slow start
 - Slow-start restart after an idle period

TCP congestion control: CWND graph




Implementation details

Update Time: Per RTT vs Per ACK

Implementation details

Update Time: Per RTT vs Per ACK

Double CWND per round-trip time




Implementation details

Update Time: Per RTT vs Per ACK

Double CWND per round-trip time

Implemented as $CWND = CWND + 1$ (per ACK)




Implementation details

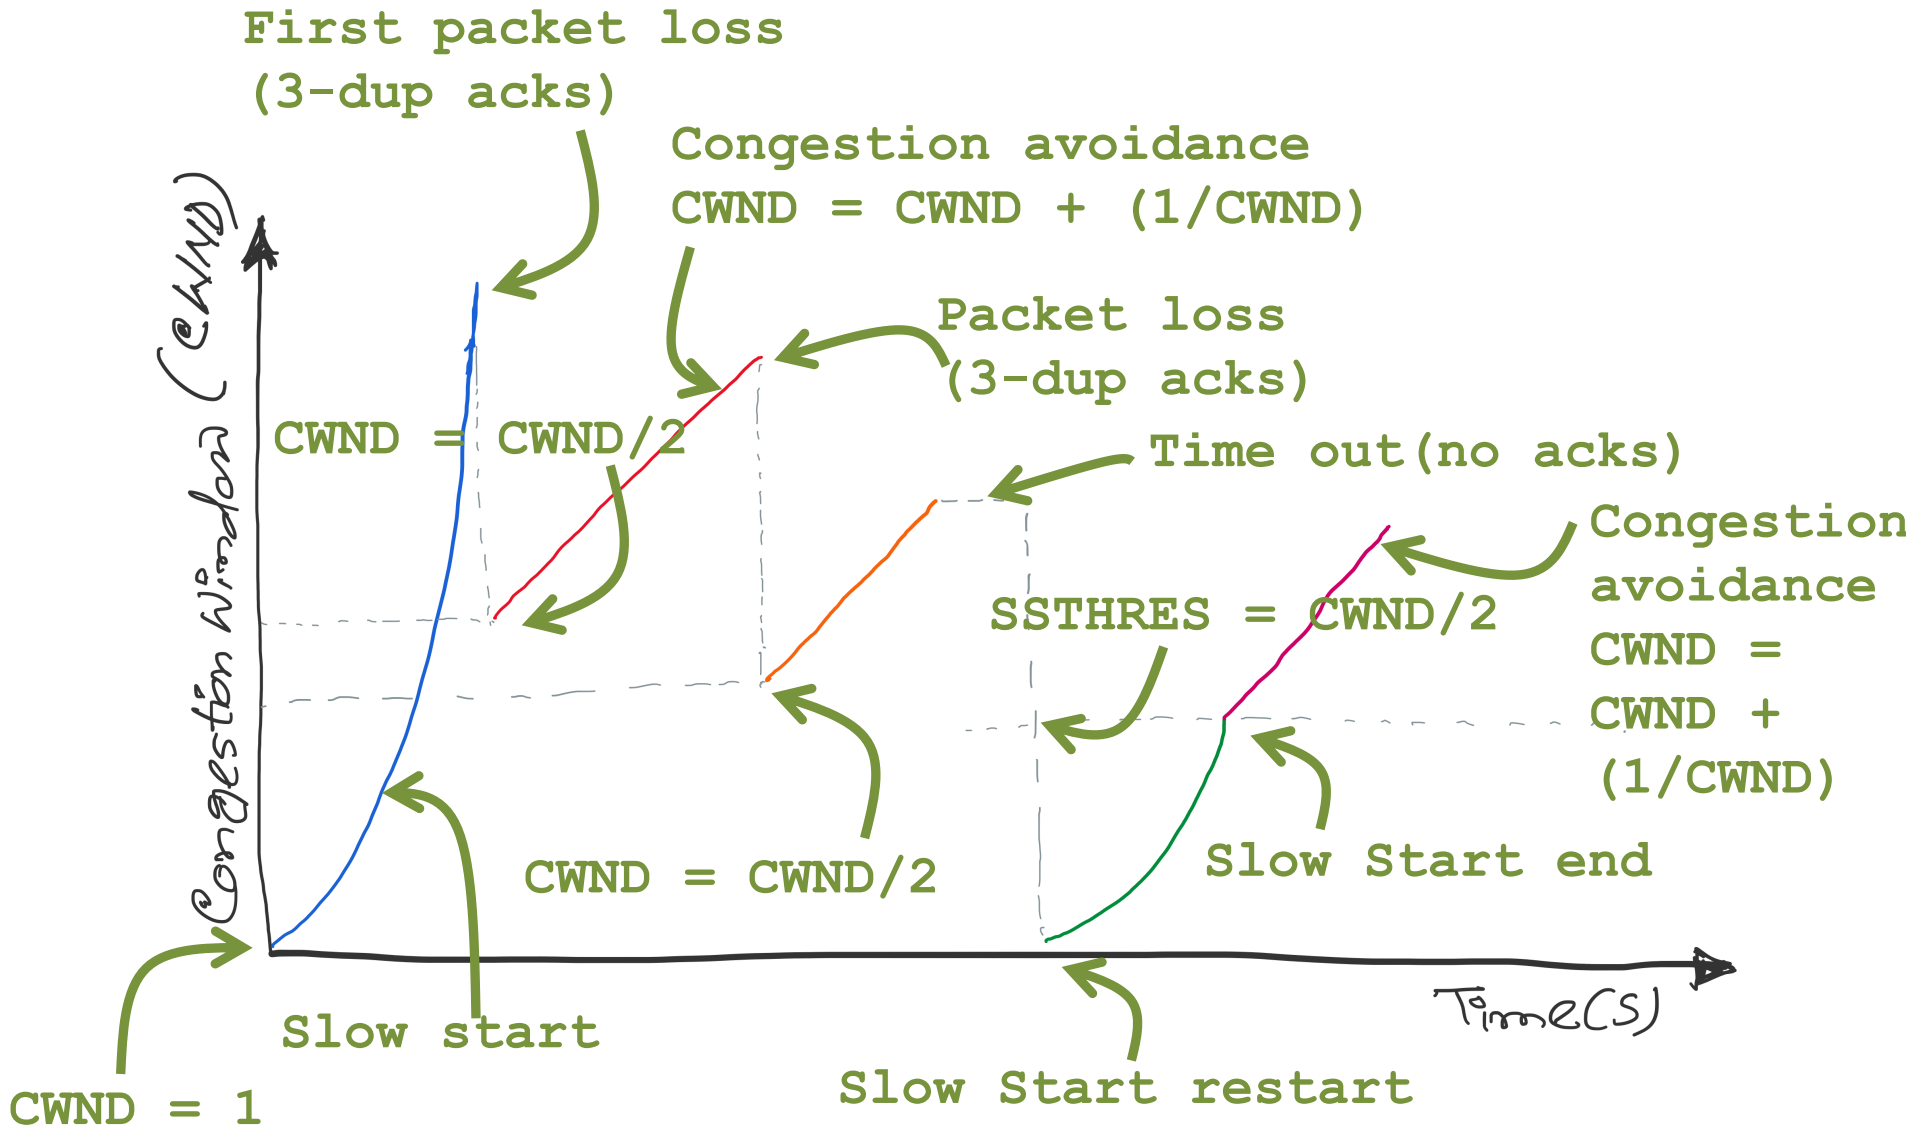
Update Time: Per RTT vs Per ACK

Linear increase of CWND per round-trip time

Implemented as $CWND = CWND + 1/CWND$ (per ACK)

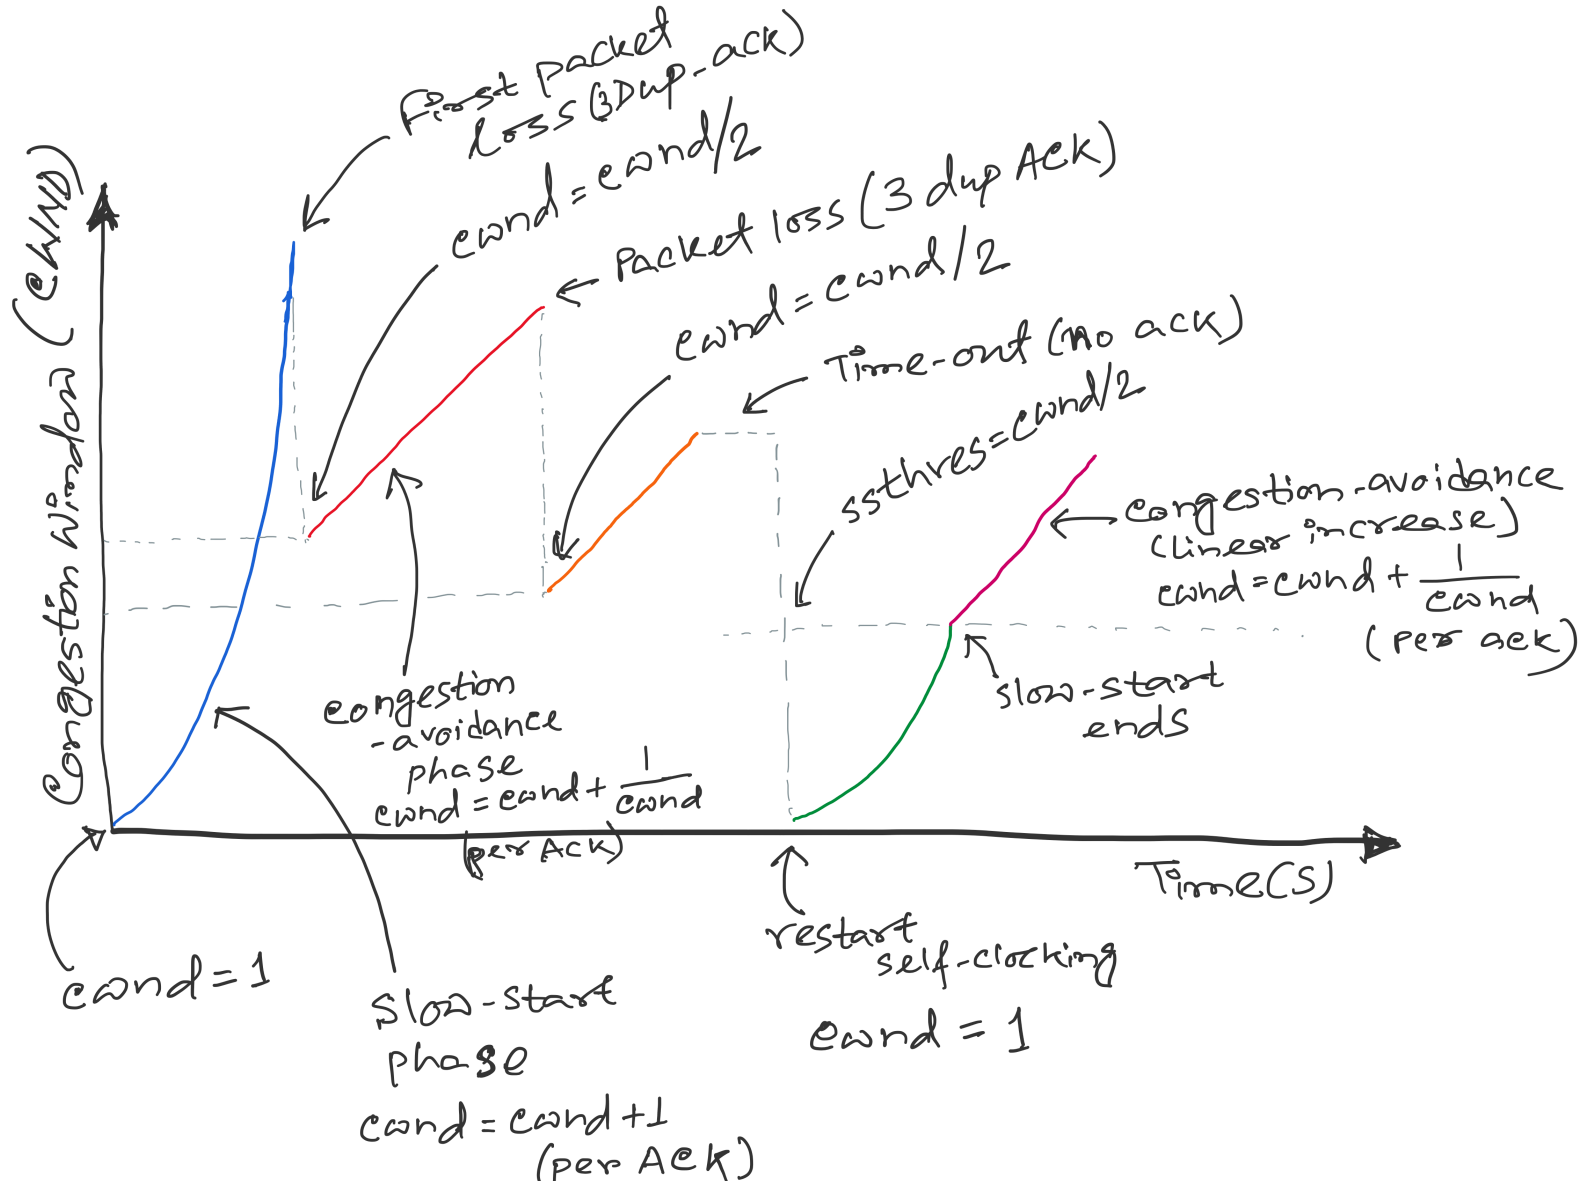


TCP congestion control: CWND graph



TCP congestion control

SSTHRES = Slow-start threshold




1. Development of Congestion Control over TCP versions

TCP versions

- TCP is not perfectly homogenous (200+)


4.2 BSD first widely available release of TCP/IP (1983)



TCP Congestion Controls

- Tahoe (Jacobson 1988)
 - Slow Start
 - Congestion Avoidance
 - Fast Retransmit (FR)

TCP Tahoe (Jacobson 1988)



SS: Slow Start

CA: Congestion Avoidance

Summary: Tahoe

- Basic ideas


- Gently probe network for spare capacity
- Drastically reduce rate on congestion
- Windowing: self-clocking
- Other functions: round trip time estimation, error recovery

```
for every ACK {  
    if ( $w < ssthresh$ ) then  $w++$       (SS)  
    else  $w += 1/w$                   (CA)  
}  
for every loss {  
     $ssthresh = w/2$   
     $w = 1$   
}
```


TCP Congestion Controls

- Tahoe (Jacobson 1988)
 - Slow Start
 - Congestion Avoidance
 - Fast Retransmit (FR)
- Reno (Jacobson 1990)
 - Fast Recovery


Tahoe's CWND recovers slowly



SS: Slow Start

CA: Congestion Avoidance

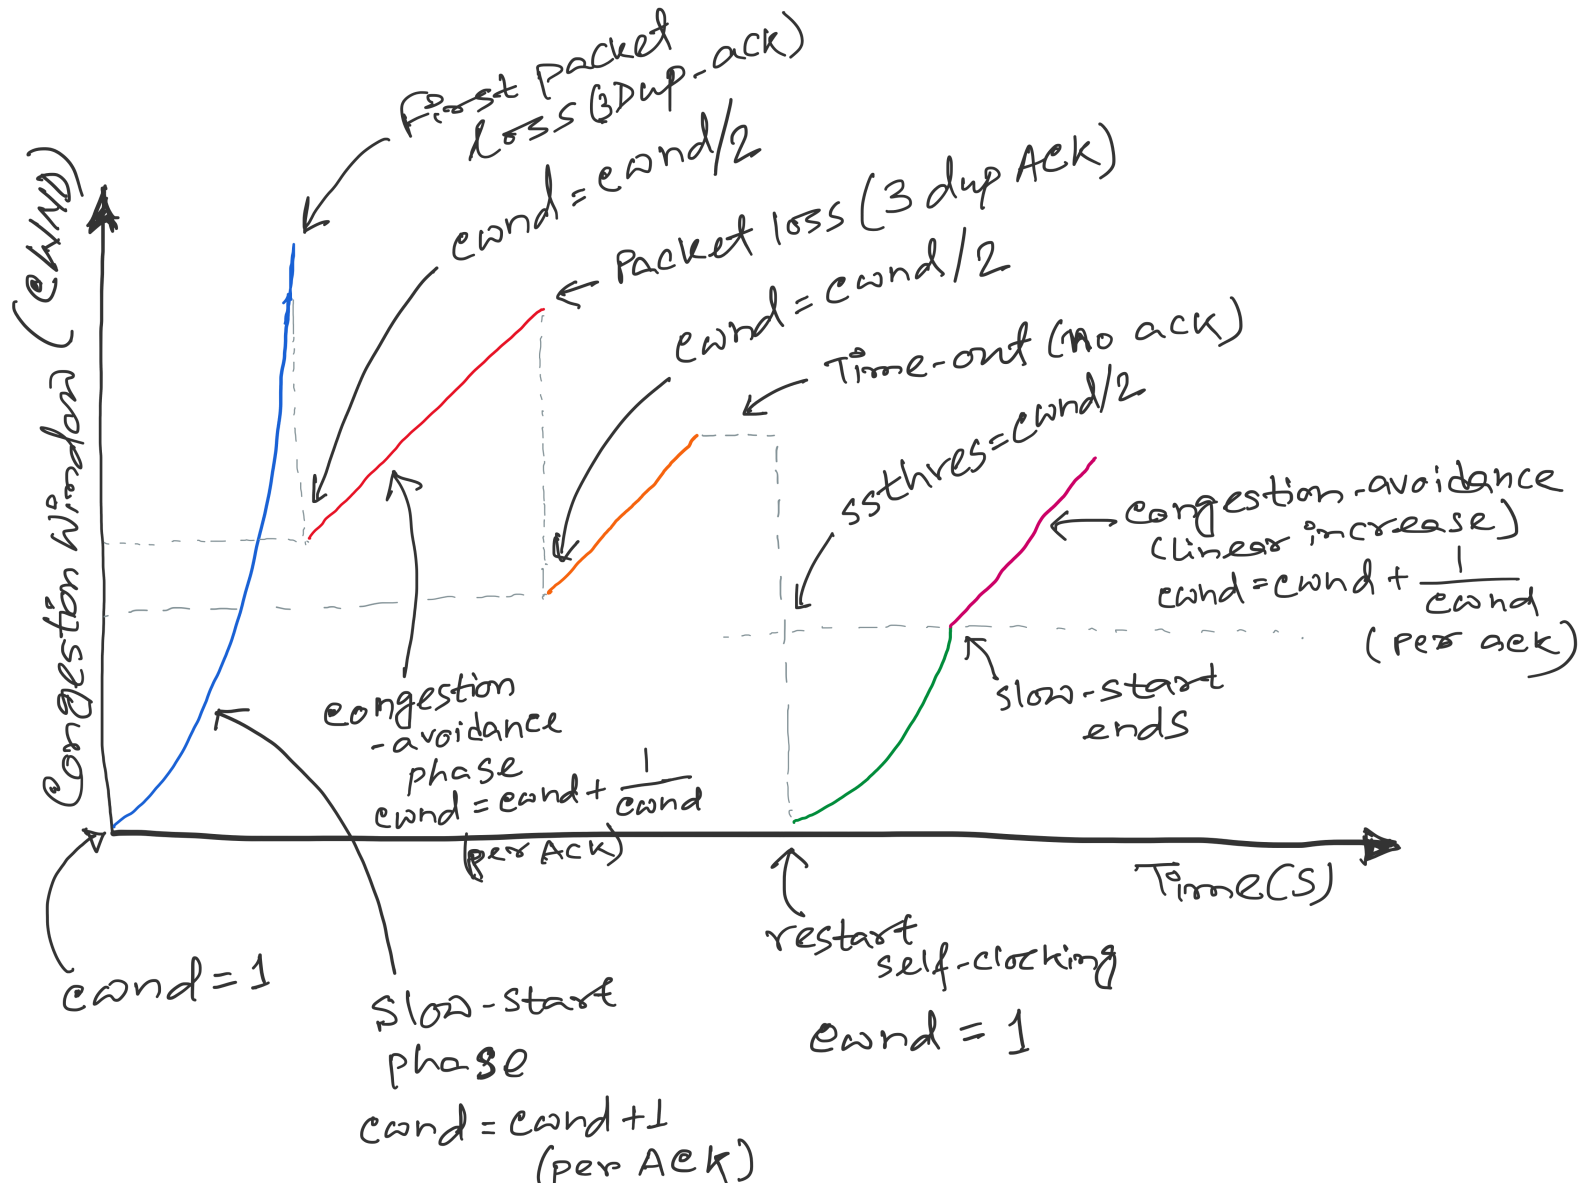
TCP Reno




SS: Slow Start
CA: Congestion Avoidance

Fast retransmission/Fast recovery

TCP Reno



TCP Reno: details of Fast Recovery



TCP Reno: details of Fast Recovery

