



Computer Networks

CMSC 417 : Spring 2024



Topics:

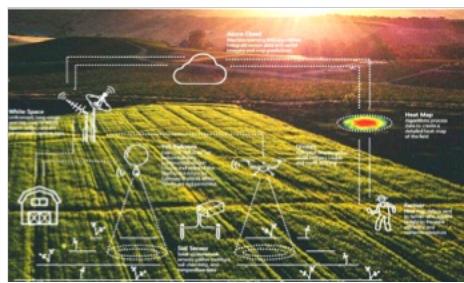
- A) TCP details (Textbook chapter 6)
- B) Link layer: Introduction, Ethernet (Textbook chapter 2)

Nirupam Roy

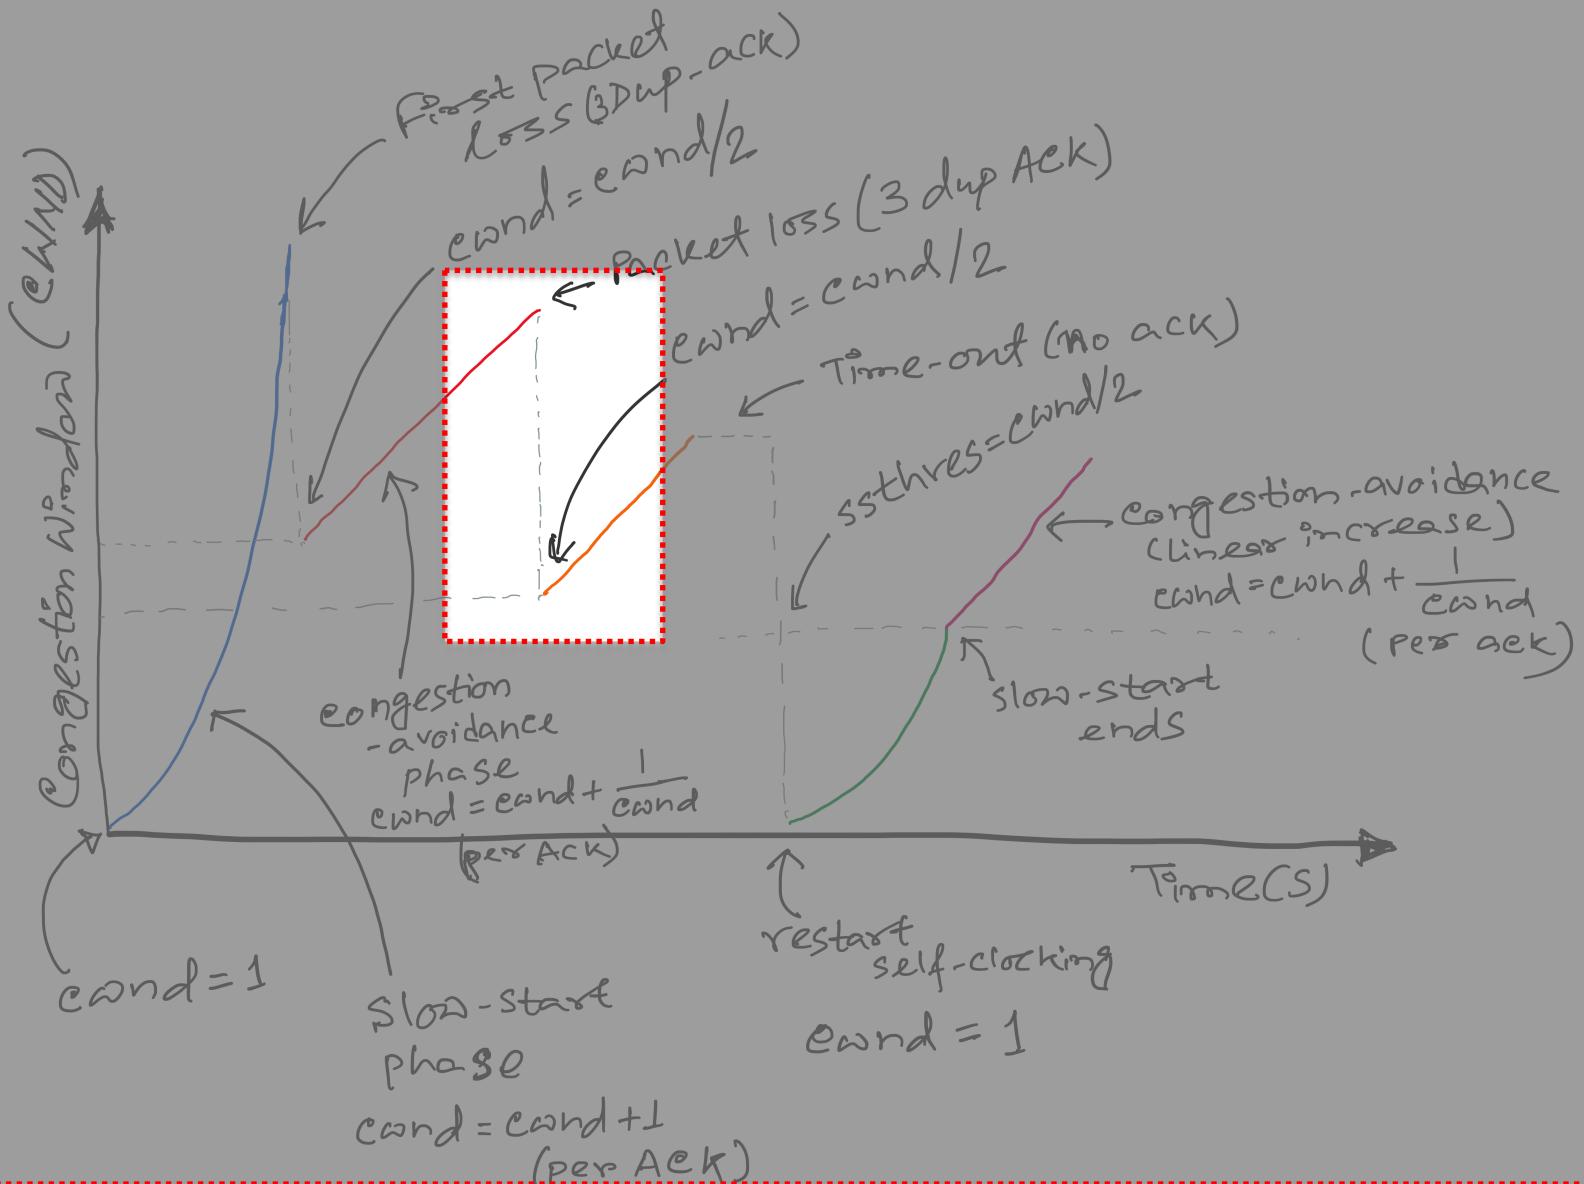
Tu-Th 2:00-3:15pm

CSI 2117

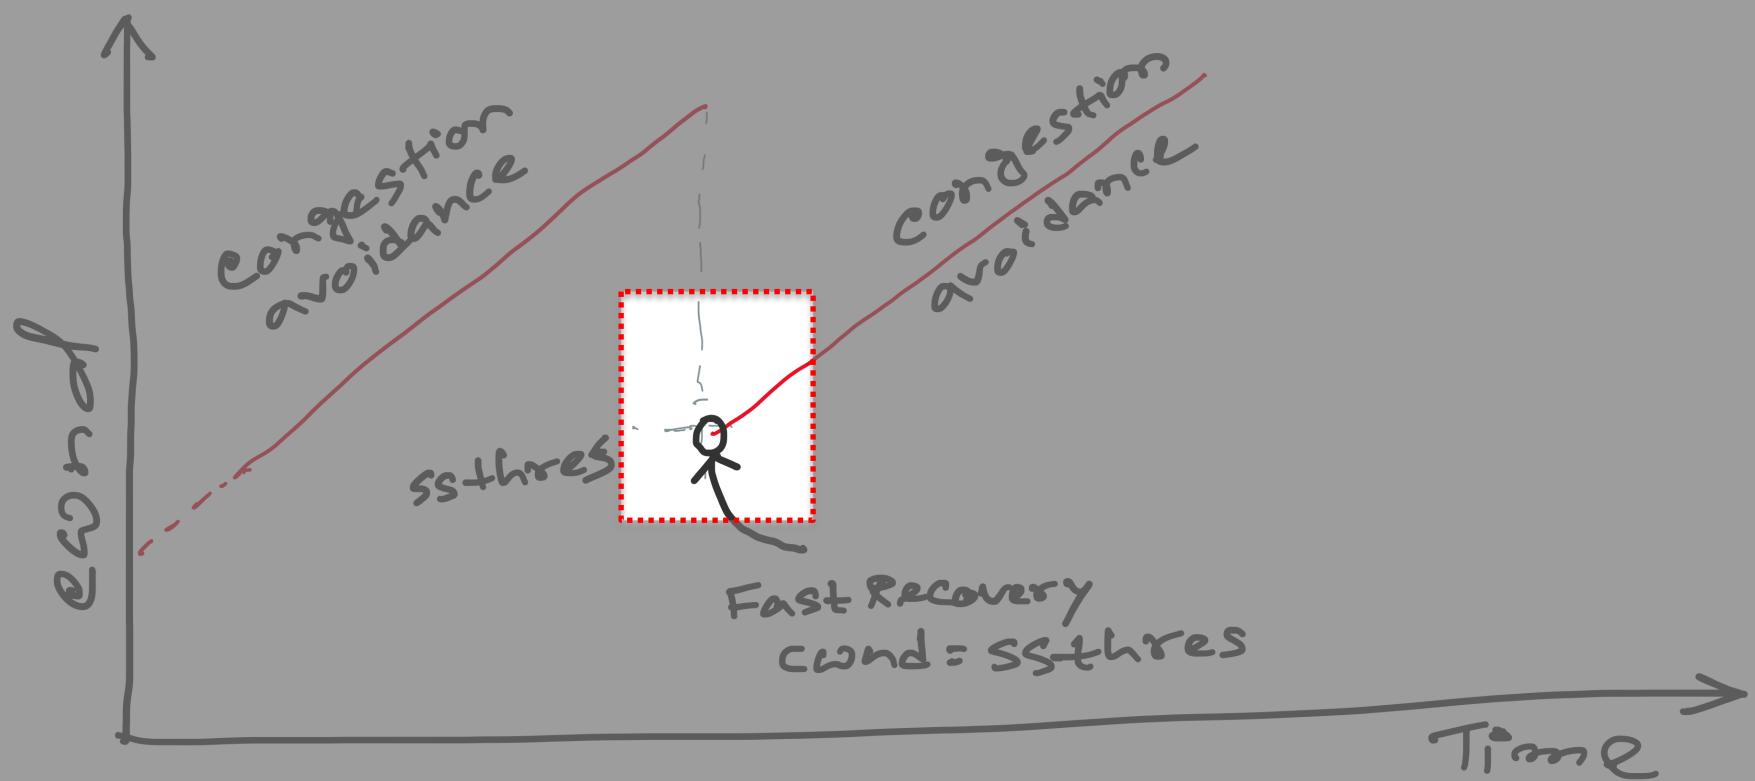
April 2nd, 2024



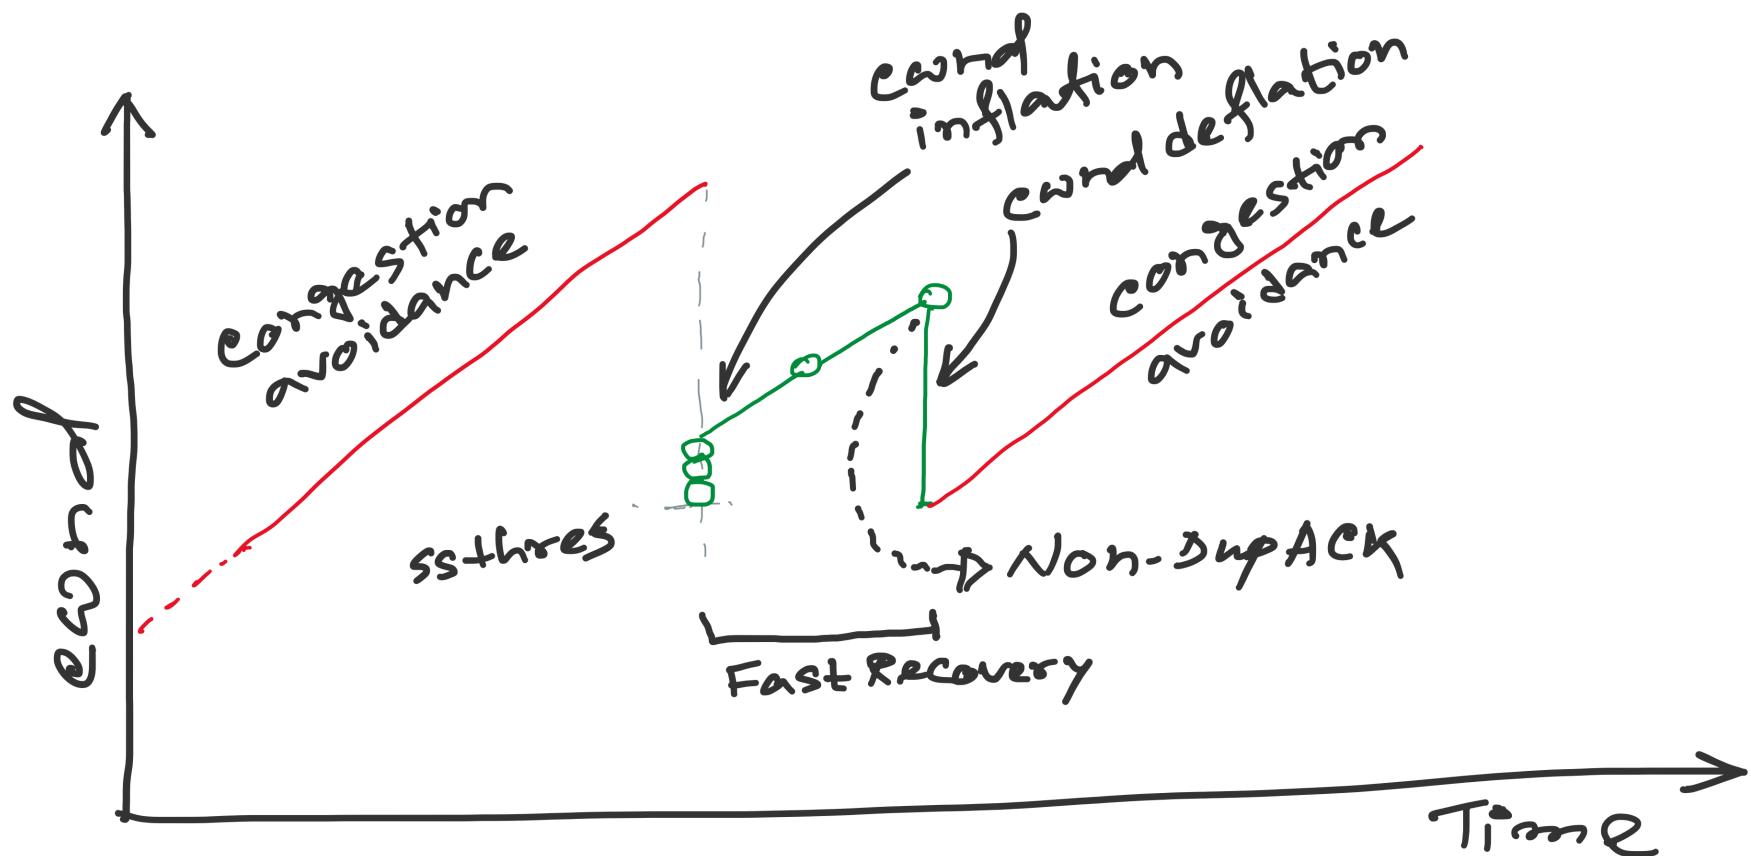
TCP Reno



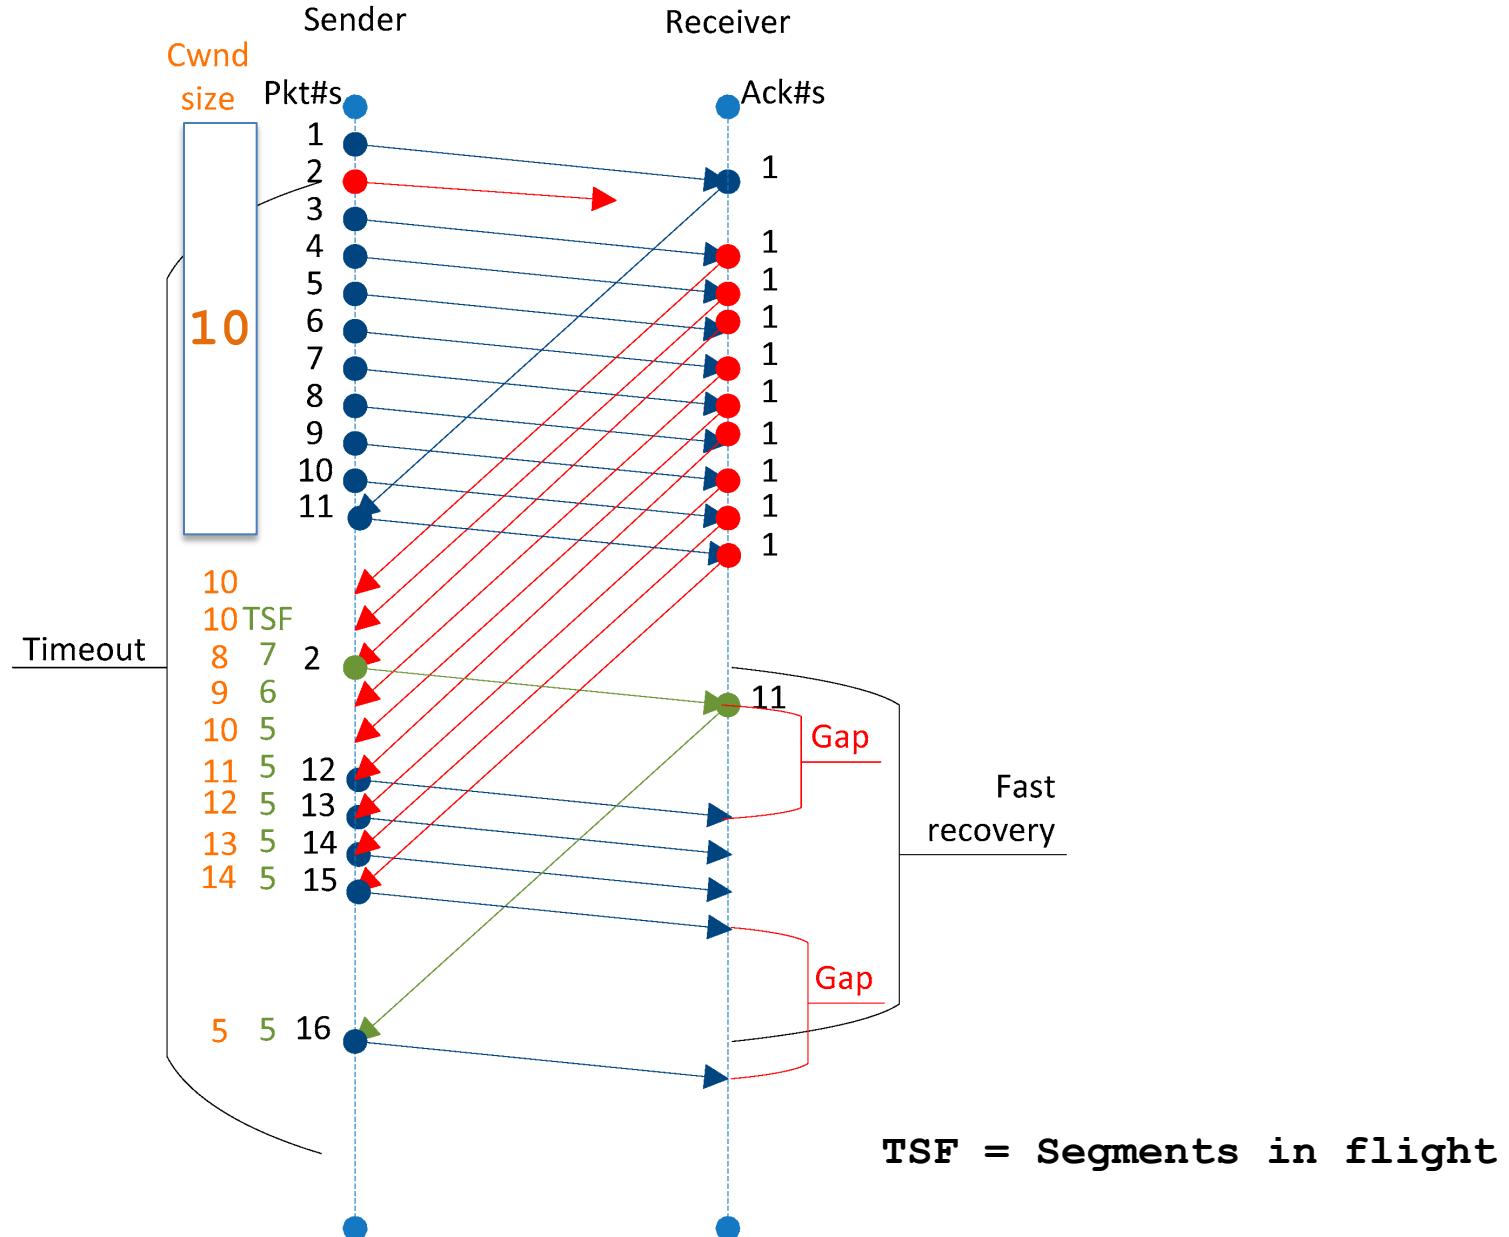
TCP Reno: details of Fast Recovery



TCP Reno: details of Fast Recovery



Fast Retransmit/Fast Recovery (FR/FR)

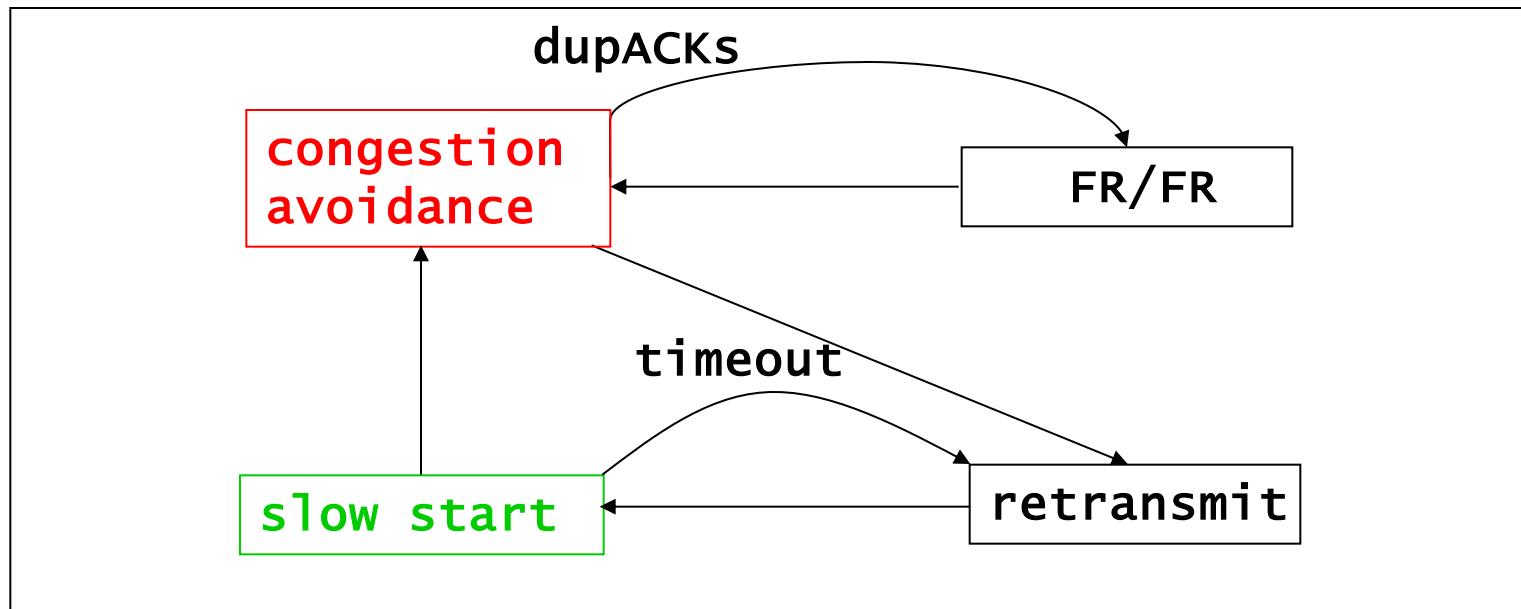


Fast Retransmit/Fast Recovery (FR/FR)

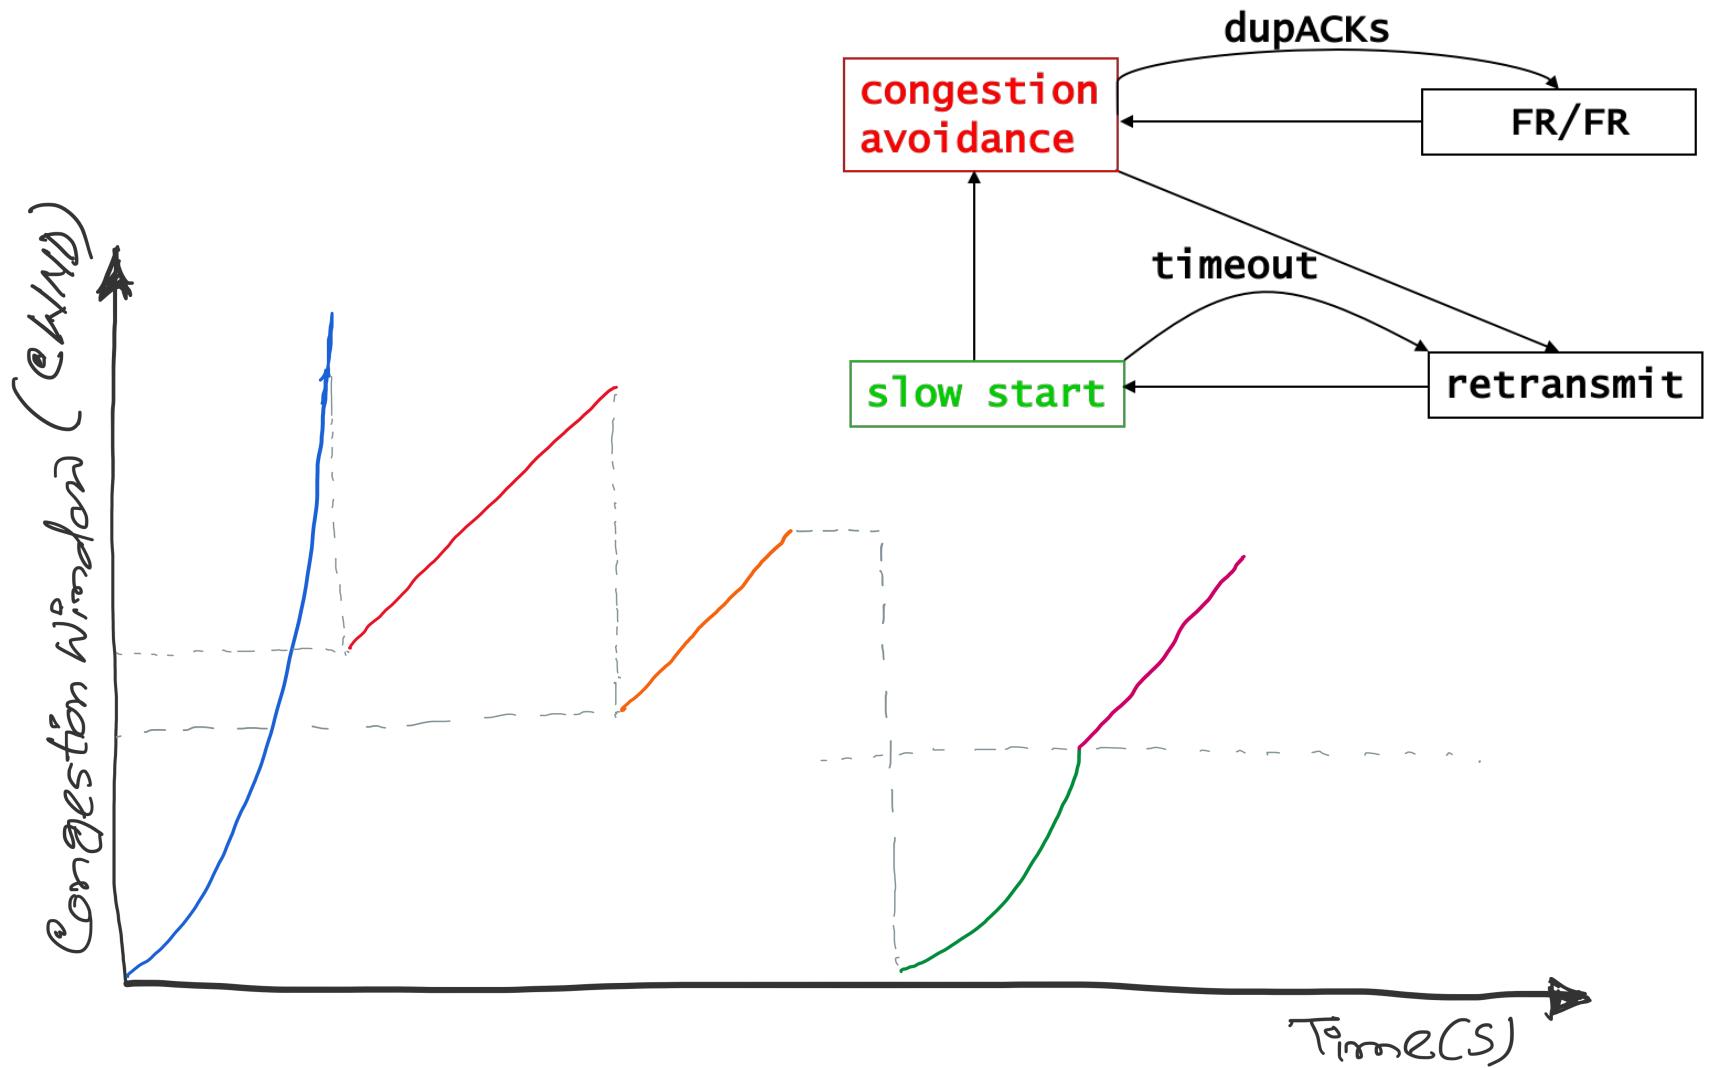
- Motivation: prevent ‘pipe’ from emptying after fast retransmit
- Idea: each dupACK represents a packet having left the pipe (successfully received)
- Enter FR/FR after 3 dupACKs
 - Set $ssthresh \leftarrow \max(\text{flightsize}/2, 2)$
 - Retransmit lost packet
 - Set $cwnd \leftarrow ssthresh + ndup$ (window inflation)
 - Wait till $W = \min(\text{awnd}, cwnd)$ is large enough; transmit new packet(s)
 - On non-dup ACK (1 RTT later), set $cwnd \leftarrow ssthresh$ (window deflation)
- Enter CA

Summary: Reno

- Basic ideas
 - Fast recovery avoids slow start
 - dupACKs: fast retransmit + fast recovery
 - Timeout: fast retransmit + slow start

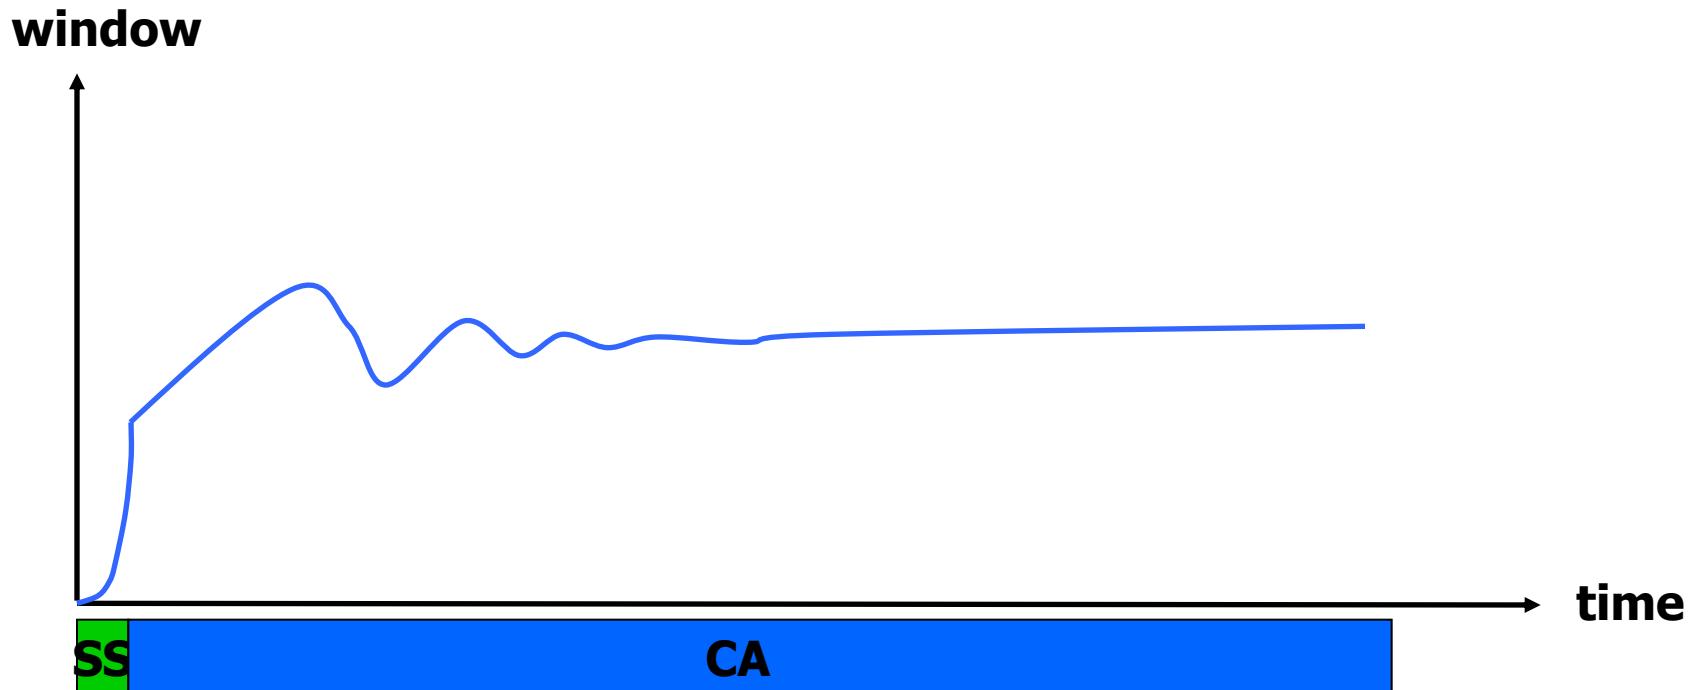


TCP congestion control: CWND graph



TCP Vegas

(Brakmo & Peterson 1994)

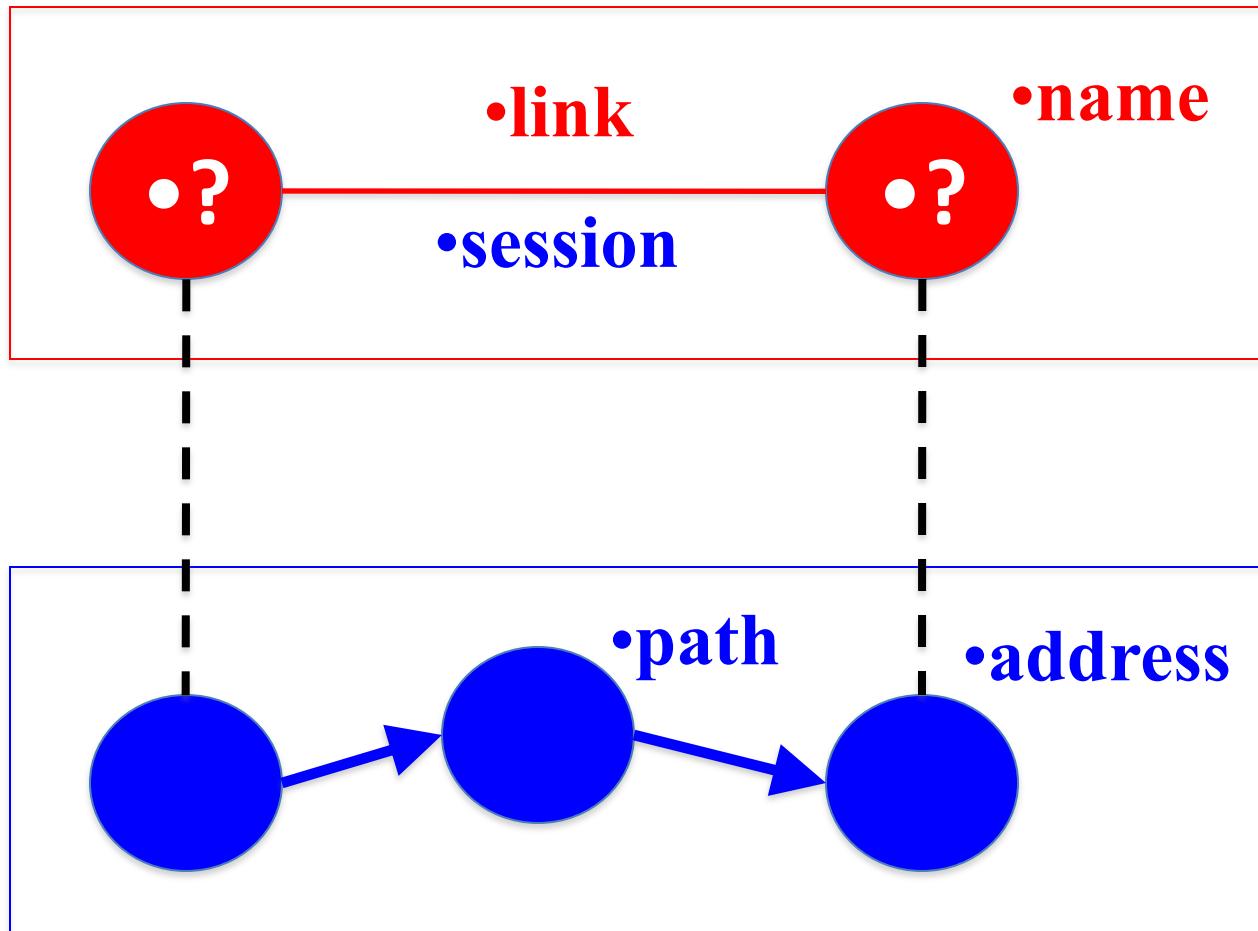


- Considers packet delay, not loss, as a hint for congestion
- Reno with a new congestion avoidance algorithm
- Converges (provided buffer is large) !

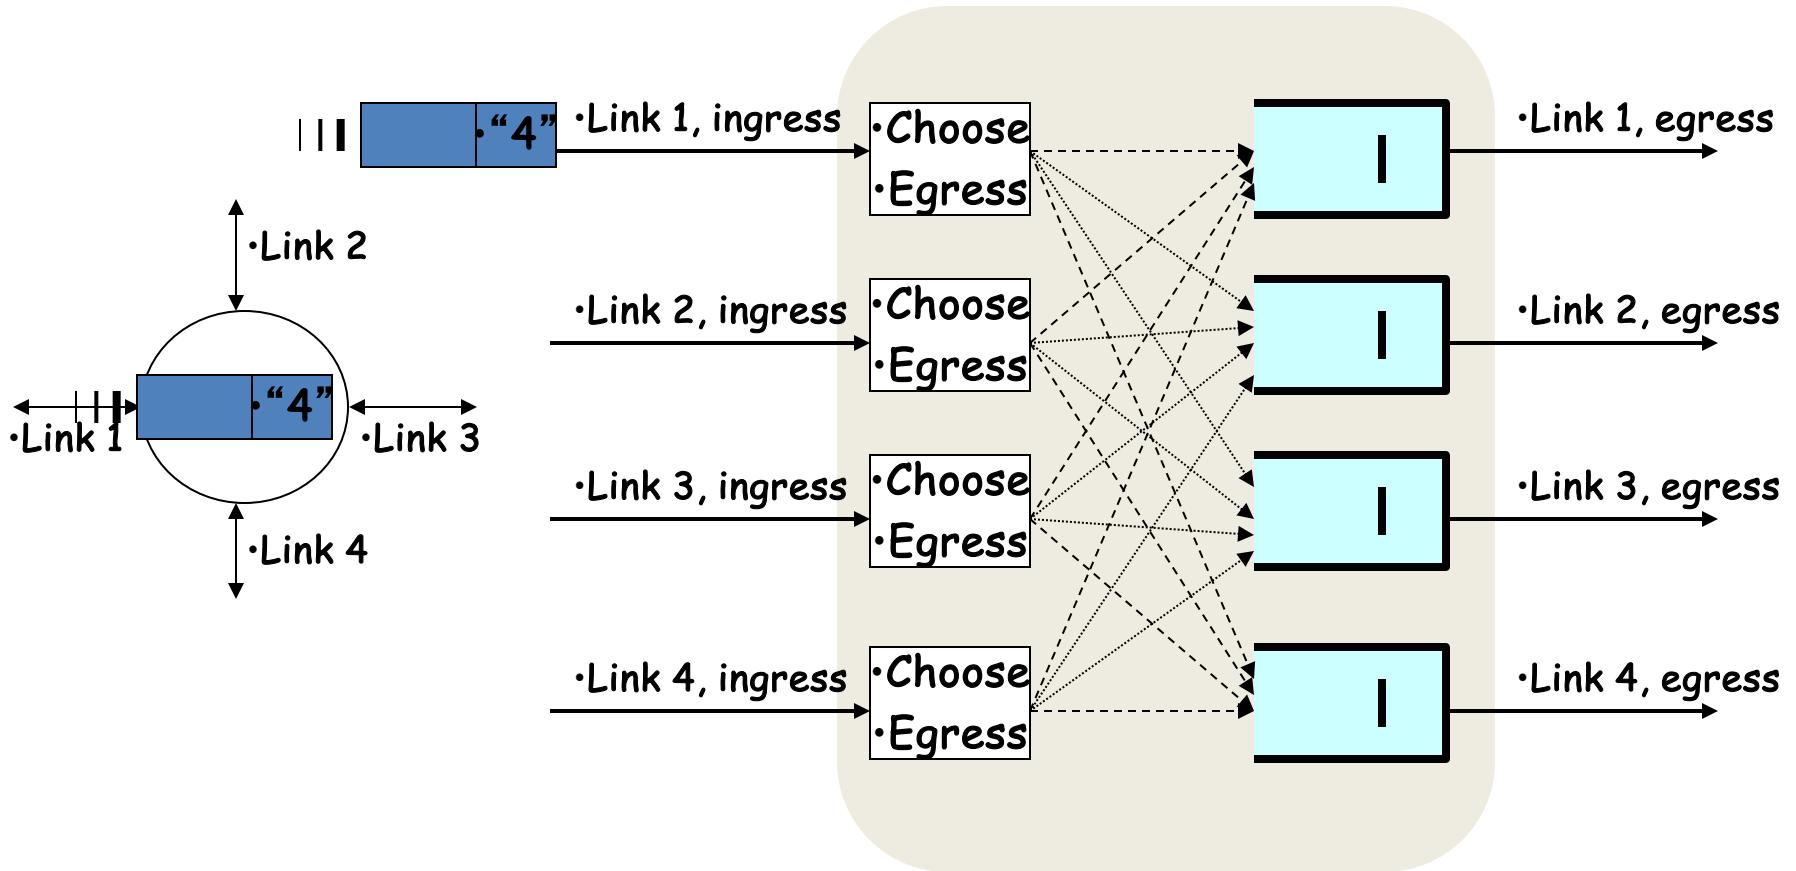
2. Router-assisted Congestion Control

Congestion Control discussed so far

- What can the *end-points* do to collectively to make good use of shared underlying resources?



Packet Switching and Forwarding



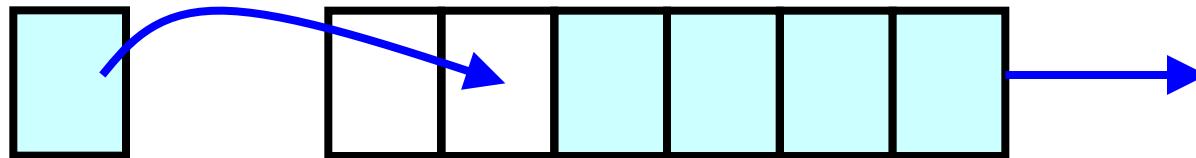
Queue Management Issues

- Scheduling discipline
 - Which packet to send?
 - Some notion of fairness? Priority?
- Drop policy
 - When should you discard a packet?
 - Which packet to discard?
- Goal: balance throughput and delay
 - Huge buffers minimize drops, but add to queuing delay (thus higher RTT, longer slow start, ...)

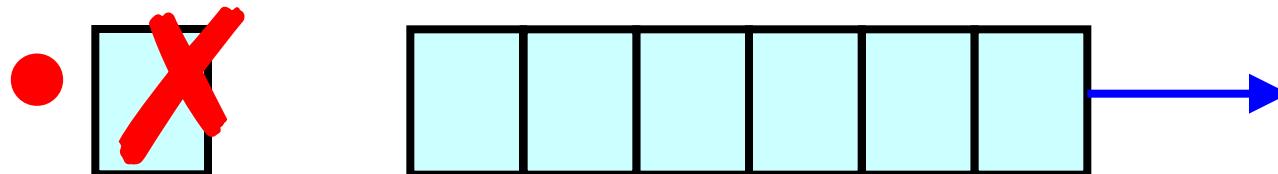
Queue Management 1: Drop Policy

FIFO Scheduling and Drop-Tail

- Access to the bandwidth: first-in first-out queue
 - Packets only differentiated when they arrive

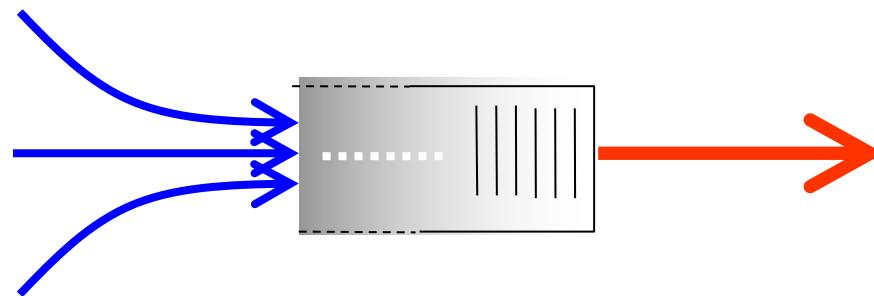


- Access to the buffer space: drop-tail queuing
 - If the queue is full, drop the incoming packet



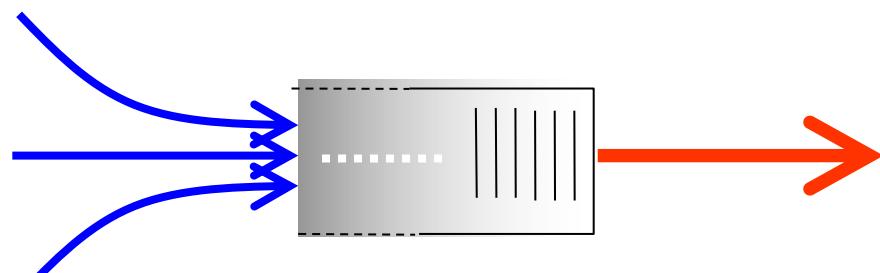
Bursty Loss From Drop-Tail Queuing

- TCP depends on packet loss
 - Packet loss is indication of congestion
 - TCP additive increase drives network into loss
- Drop-tail leads to *bursty* loss
 - Congested link: many packets encounter full queue
 - Synchronization: many connections lose packets at once



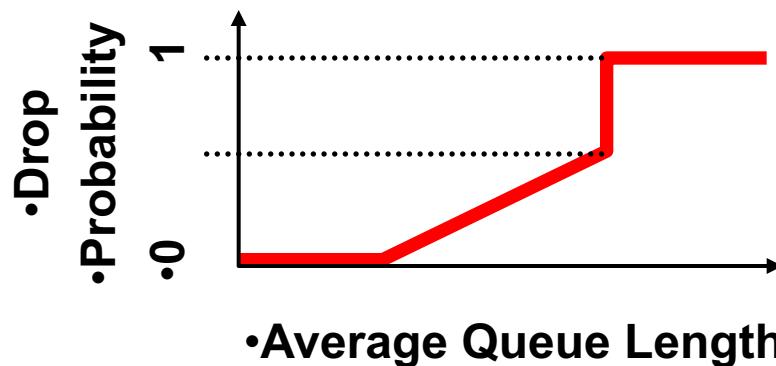
Slow Feedback from Drop Tail

- Feedback comes when buffer is completely full
 - ... even though the buffer has been filling for a while
- Plus, the filling buffer is increasing RTT
 - ... making detection even slower
- Better to give early feedback
 - Get 1-2 connections to slow down before it's too late!



Random Early Detection (RED)

- Router notices that queue is getting full
 - ... and randomly drops packets to signal congestion
- Packet drop probability
 - Drop probability increases as queue length increases
 - Else, set drop probability $f(\text{avg queue length})$



Properties of RED

- **Drops packets before queue is full**
 - In the hope of reducing the rates of some flows
- **Drops packet in proportion to each flow's rate**
 - High-rate flows selected more often
- **Drops are spaced out in time**
 - Helps desynchronize the TCP senders
- **Tolerant of burstiness in the traffic**
 - By basing the decisions on average queue length

Problems With RED

- Hard to get tunable parameters just right
 - How early to start dropping packets?
 - What slope for increase in drop probability?
 - What time scale for averaging queue length?
- RED has mixed adoption in practice
 - If parameters aren't set right, RED doesn't help
- Many other variations in research community
 - Names like “Blue” (self-tuning), “FRED”...

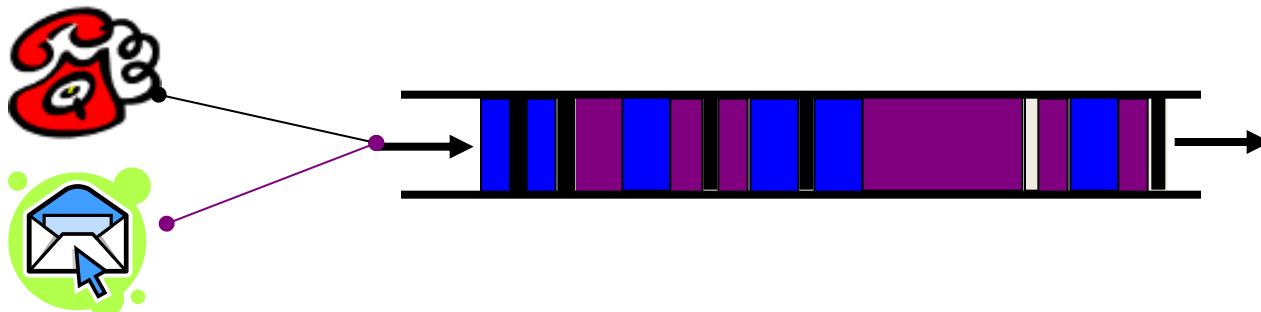
Feedback: From Loss to Notification

- Early dropping of packets
 - Good: gives early feedback
 - Bad: has to drop the packet to give the feedback
- Explicit Congestion Notification (ECN)
 - Router marks the packet with an ECN bit
 - Sending host interprets as a sign of congestion
 - Requires participation of hosts and the routers

Queue Management 2: Link Scheduling

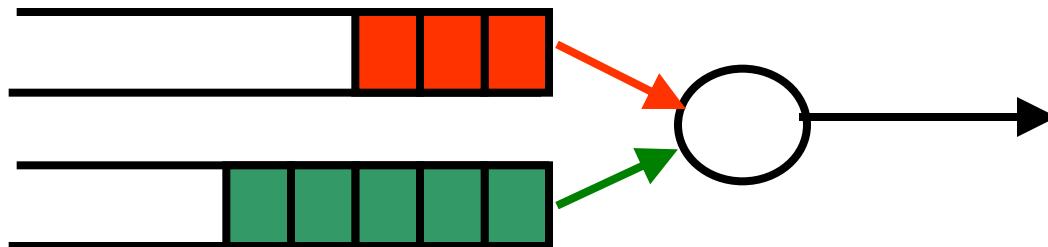
First-In First-Out Scheduling

- First-in first-out scheduling
 - Simple, but restrictive
- Example: two kinds of traffic
 - Voice over IP needs low delay
 - E-mail is not that sensitive about delay
- Voice traffic waits behind e-mail



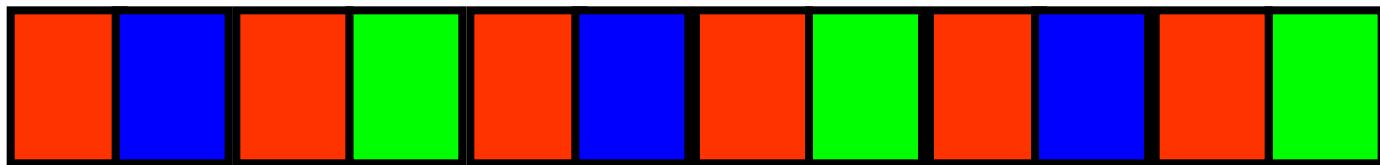
Strict Priority

- Multiple levels of priority
 - Always transmit high-priority traffic, when present
- Isolation for the high-priority traffic
 - Almost like it has a dedicated link
 - Except for (small) delay for packet transmission
- But, lower priority traffic may starve ☹



Weighted Fair Scheduling

- Weighted fair scheduling
 - Assign each queue a fraction of the link bandwidth
 - Rotate across queues on a small time scale



- 50% red, 25% blue, 25% green

- Work-conserving
 - Send extra traffic from one queue if others are idle

Implementation Trade-Offs

- FIFO
 - One queue, trivial scheduler
- Strict priority
 - One queue per priority level, simple scheduler
- Weighted fair scheduling
 - One queue per class, and more complex scheduler

TCP Congestion Control

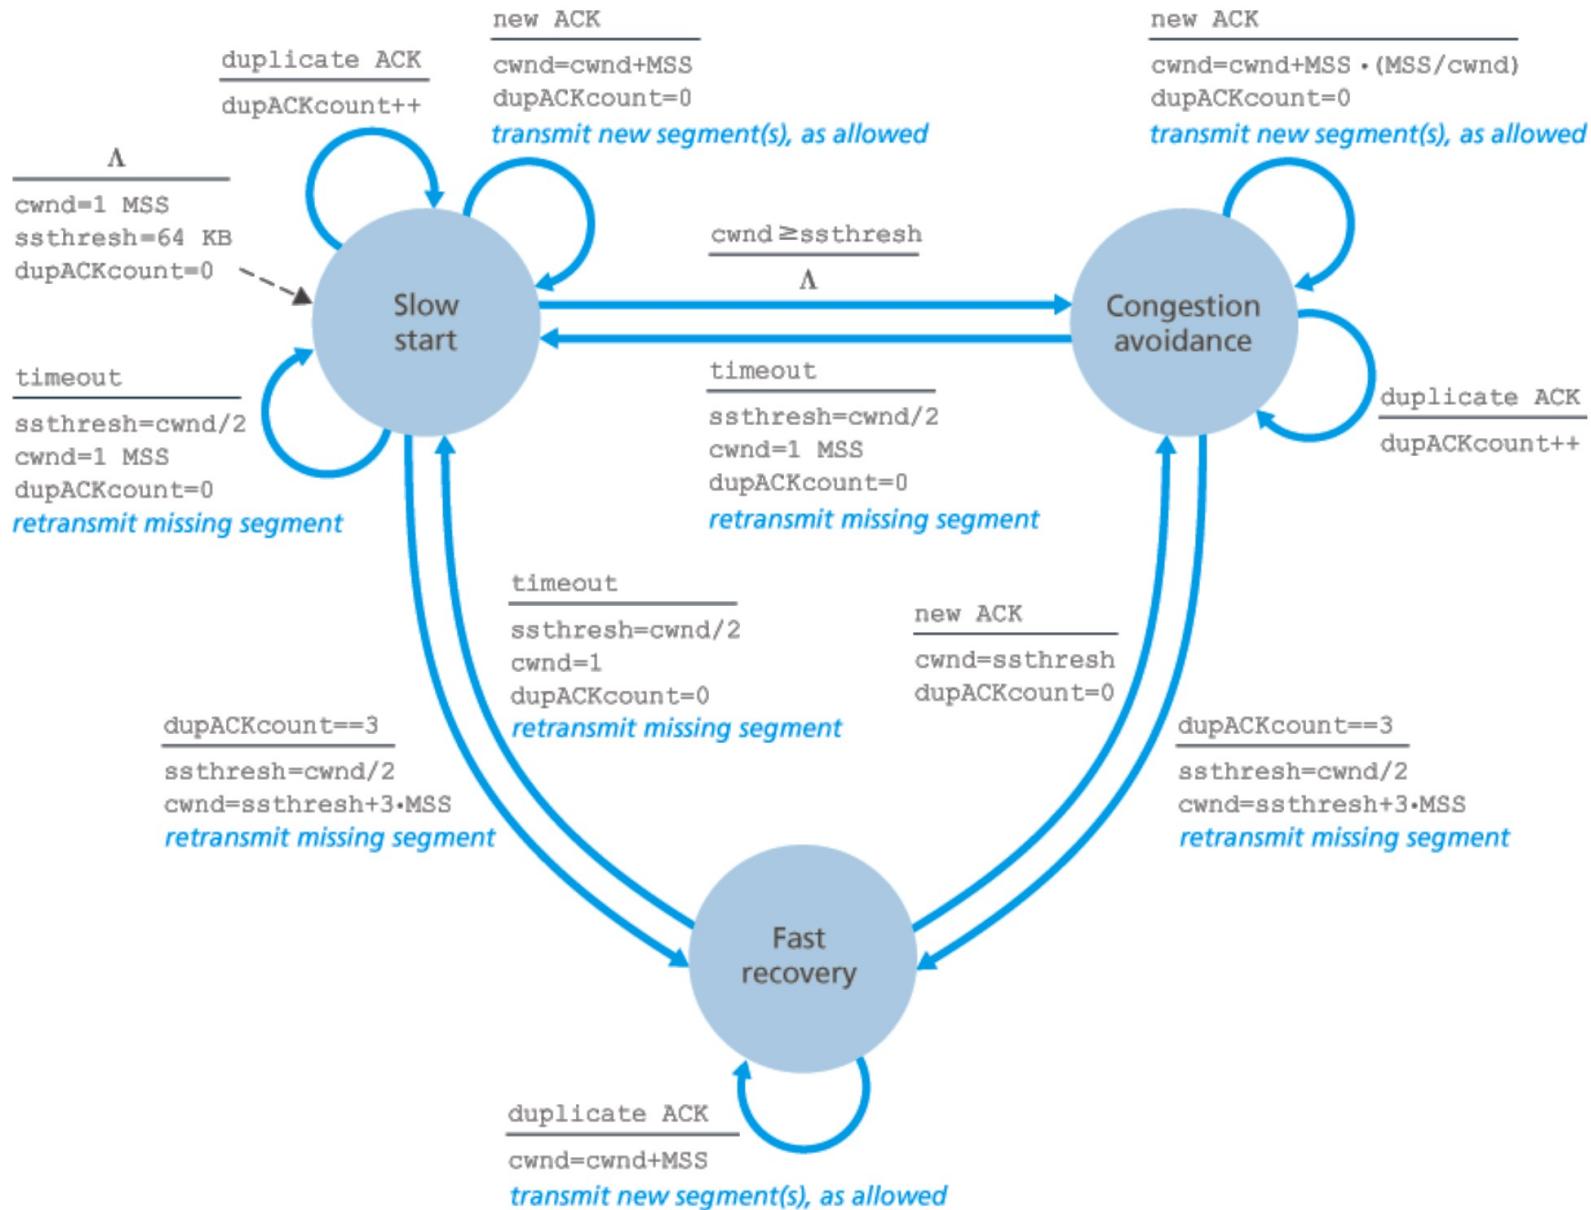
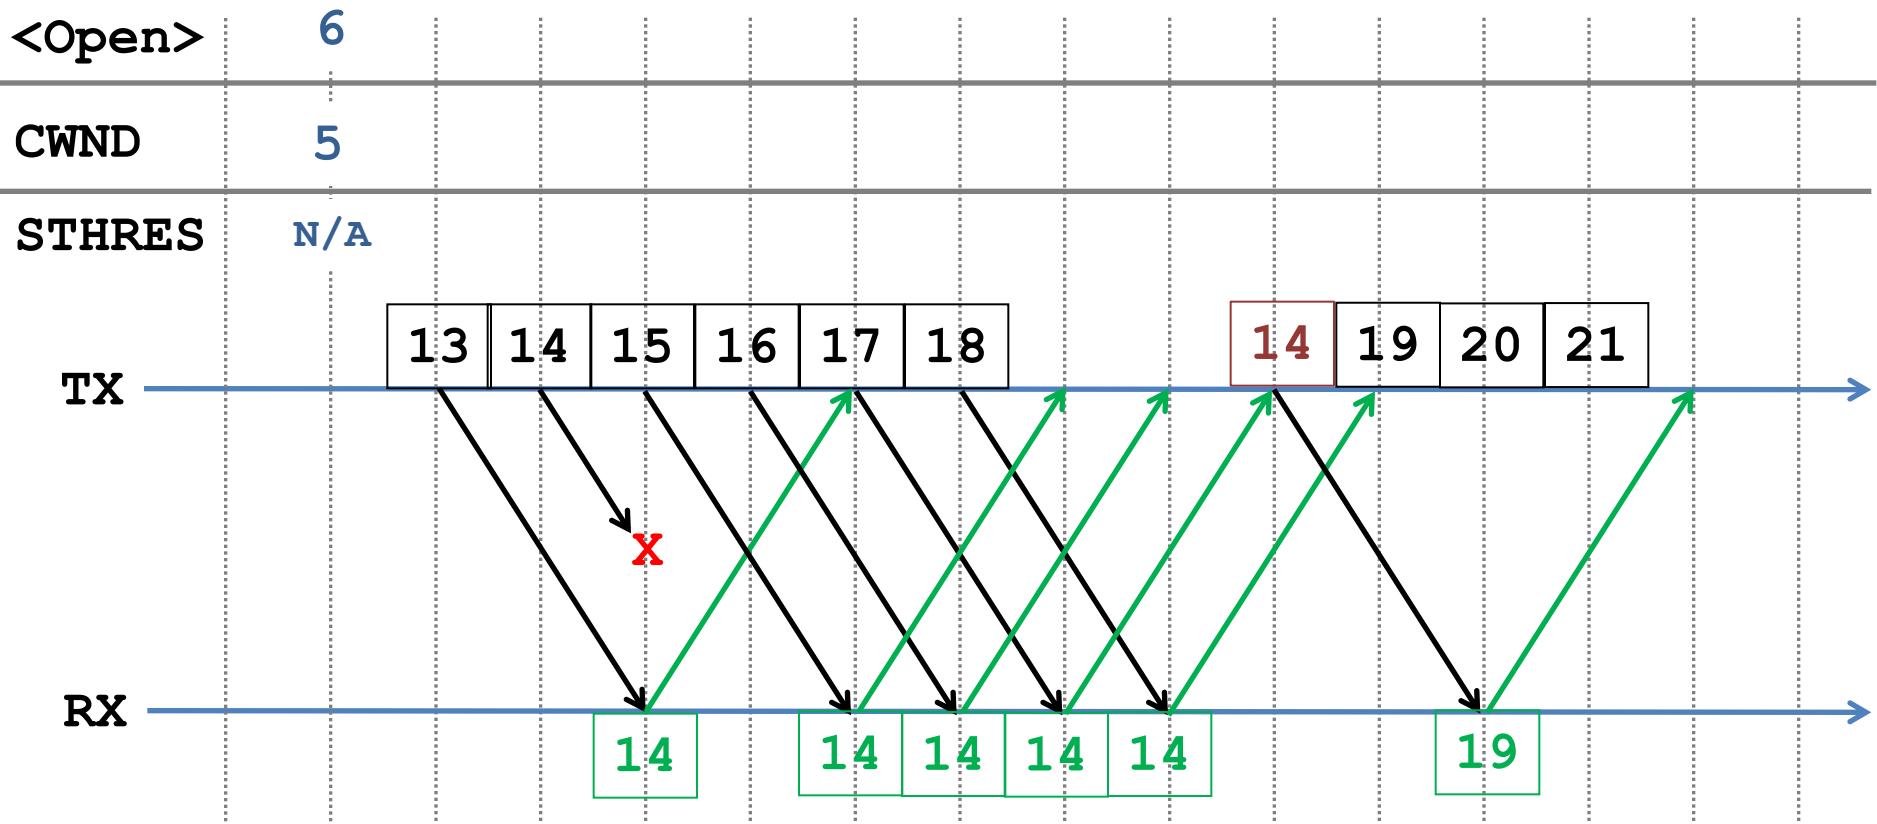
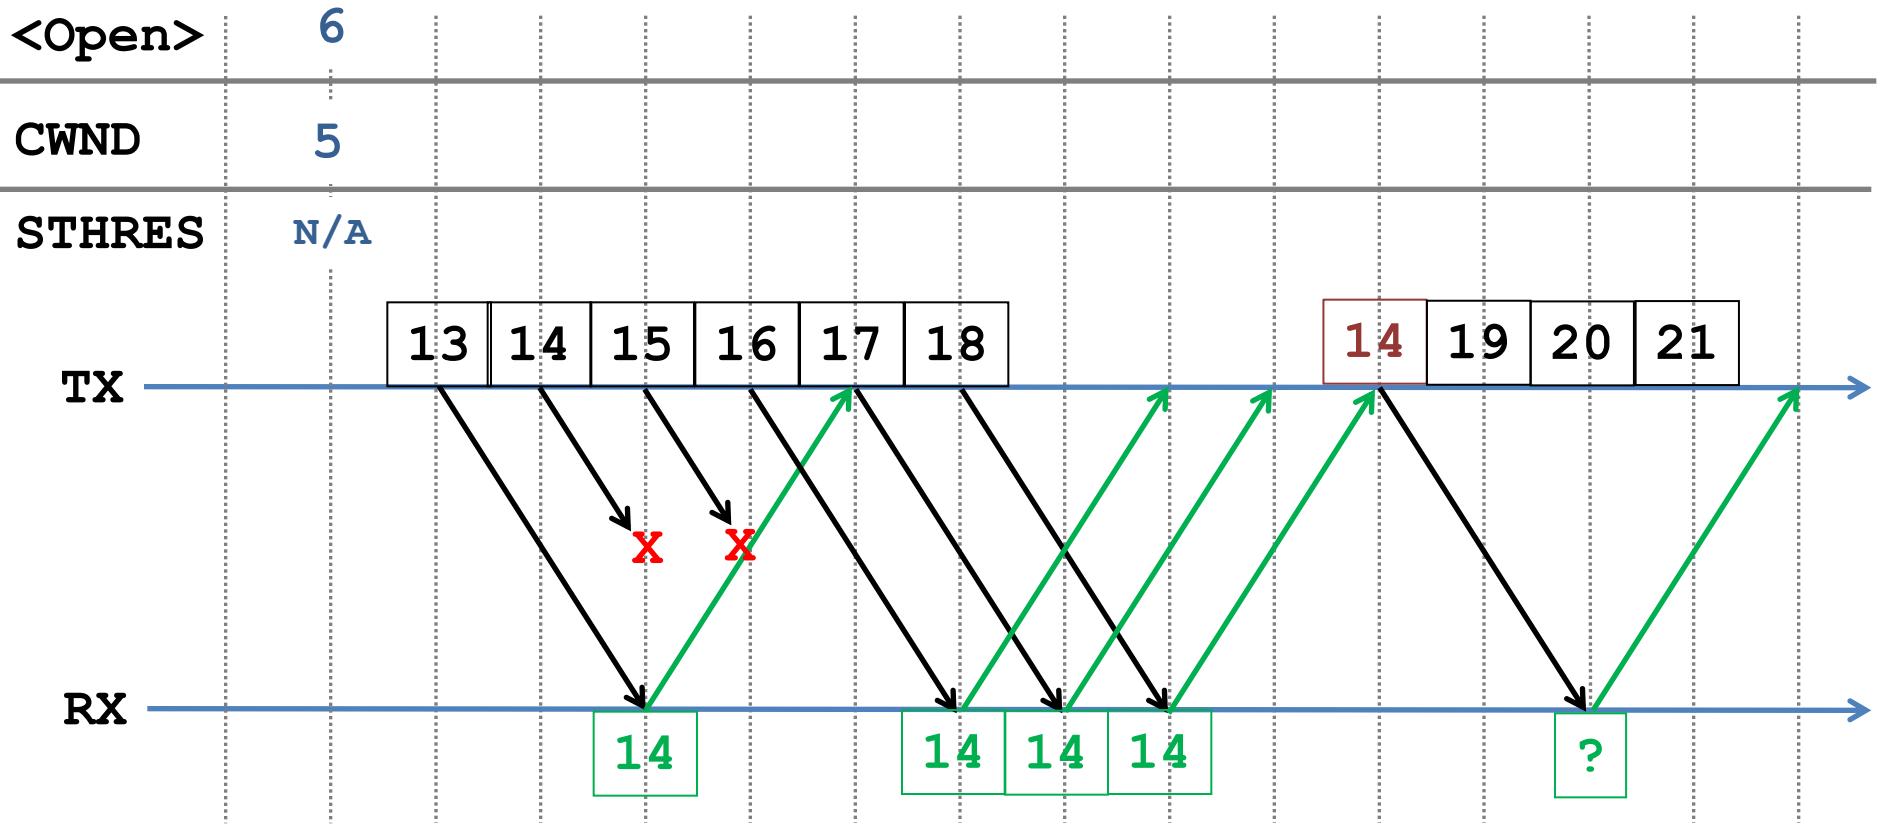


Figure 3.51 FSM description of TCP congestion control

Revisiting Dup ACKs (without Fast-Recovery)



Revisiting Dup ACKs (without Fast-Recovery)



What if there are multiple losses in an RTT?

TCP Selective ACK (SACK): RFC 2018

<https://www.rfc-editor.org/rfc/rfc2018>

TCP Sack-Permitted Option:

Kind: 4

Kind=4	Length=2
--------	----------

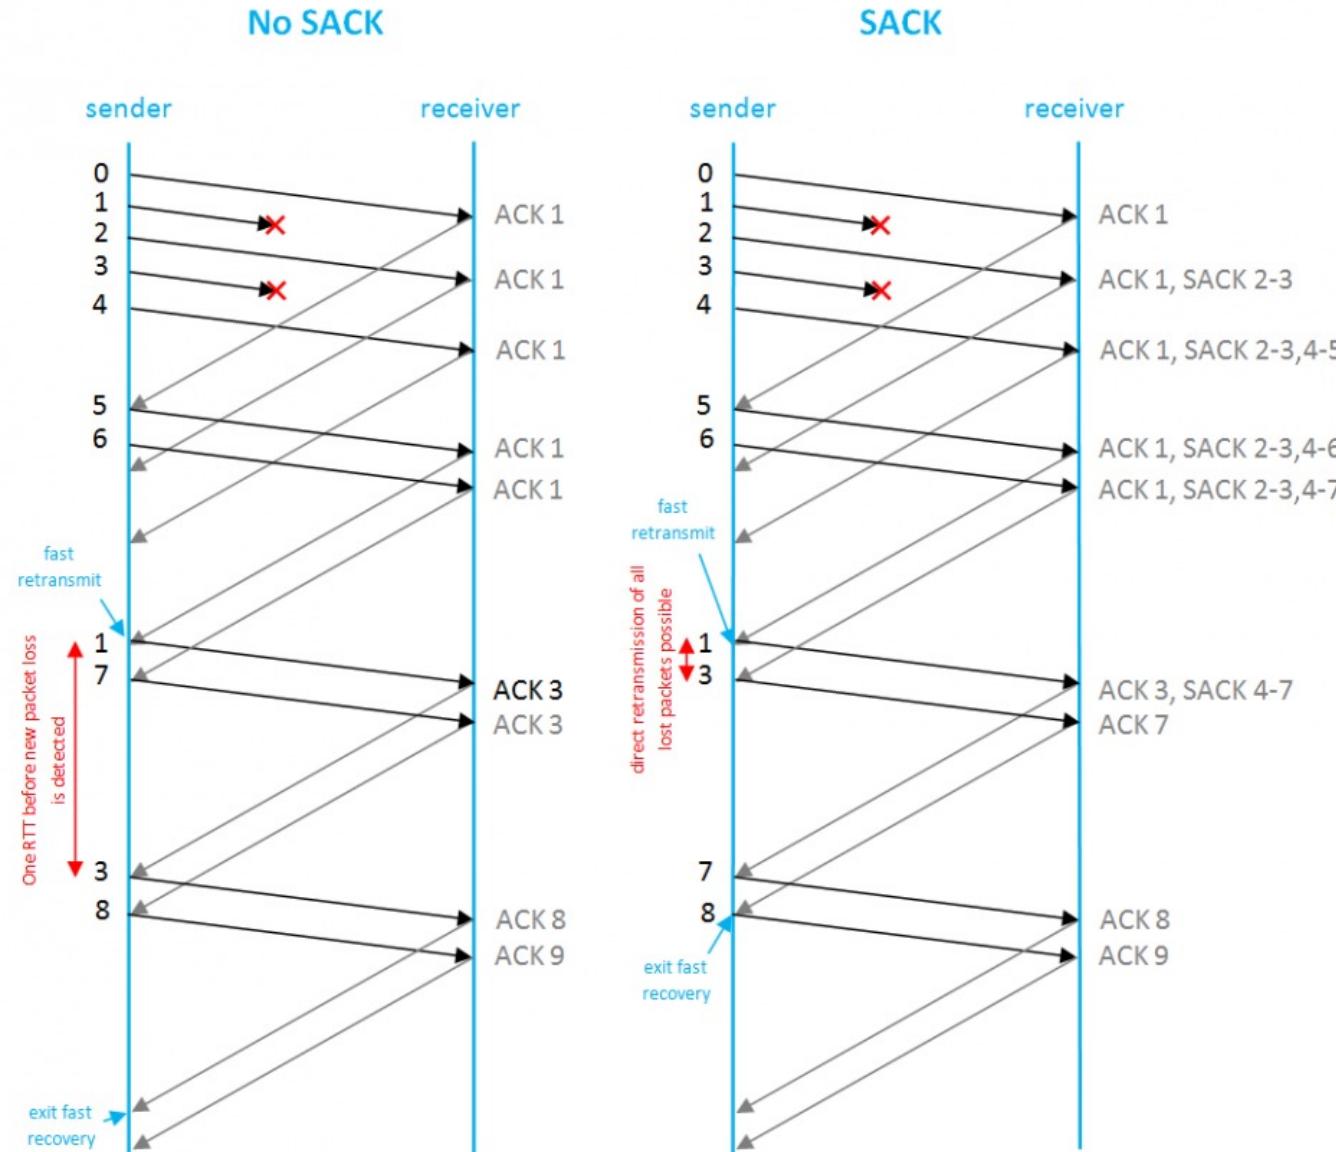
TCP SACK Option:

Kind: 5

Length: Variable

Left Edge of 1st Block	Right Edge of 1st Block
⋮	⋮
Left Edge of nth Block	Right Edge of nth Block

TCP Selective ACK (SACK): RFC 2018



[REF-1] <https://www.rfc-editor.org/rfc/rfc2018>

[REF-2] <https://www.excentis.com/blog/measuring-throughput-effect-of-used-tcp-settings/>

TCP Congestion Control with a Misbehaving Receiver

Stefan Savage, Neal Cardwell, David Wetherall, and Tom Anderson
Department of Computer Science and Engineering
University of Washington, Seattle

Slides courtesy: <https://cseweb.ucsd.edu/classes/wi01/cse222/lectures/attack-18.pdf>

The problem

- Bandwidth sharing on the Internet
 - Hosts **voluntarily** limit own data rate
 - Mechanism implemented in TCP
 - “Fair” rate determined by testing the network
 - Relies on **cooperation** between endpoints
- Why doesn’t everyone cheat?

One explanation

- Cheating requires motive **and** opportunity
- Senders (e.g., Web servers)
 - Have opportunity (could send too fast)
 - Limited motive (economic incentive to share)
- Receivers (e.g., Web clients)
 - Have competitive motive (faster Web surfing)
 - No opportunity (only receive data)... right?

What if receivers misbehave?

- A client can **implicitly** control the data rate of a remote server
 - This is not an implementation error
 - It is a weakness in the TCP specification
 - TCP's design does not consider that senders and receivers might have disjoint interests
- The vulnerability is significant...

Why the Web is faster in Seattle

